

Belief Management and Optimal Arbitration*

Benjamin Balzer[†] Johannes Schneider[‡]

May 22, 2019

Abstract

We demonstrate equivalence between the mechanism-design problem of finding the optimal arbitration mechanism and an information-design problem of managing disputing parties' beliefs in case arbitration fails to settle. We impose five axioms on our environment: ordered types, (ex-post) desirable settlement, unilateral veto rights to arbitration, exogenous rules of conflict, and ex-ante budget balance. Under these axioms any information structure in conflict maps to a unique candidate mechanism. Optimal arbitration is among these candidates. We derive an information-design problem equivalent to the arbitration problem. Its solution determines the optimal information structure for a range of objectives.

1 Introduction

Resolving conflicts through an open fight often implies high costs. Thus, attempting to resolve conflicts before they escalate to a fight is common. One of the most powerful resolution attempts is third-party arbitration. Once parties agree on arbitration, an arbitrator controls the outcome of arbitration. Still arbitrators seldom guarantee settlement and often operate *in the shadow of the fight*.

In this paper, we propose a general framework to answer the question: *How should we design arbitration?* In particular, we address how the details of the fighting stage affect optimal arbitration. This question is relevant once the arbitrator cannot guarantee settlement. If information affects equilibrium outcomes in the fighting stage, the arbitration problem is sensitive to information revelation.

*We thank Eddie Deckel, Piotr Dworczak, Willie Fuchs, Johannes Hörner, Antoine Loeper, Benny Moldovanu, Volker Nocke, Alessandro Pavan, Thomas Tröger, and Christoph Wolf for helpful comments at various stages of the paper. Some parts of this paper have earlier been circulated as part of another paper with title “Managing a Conflict.” Johannes Schneider gratefully acknowledges support from the Ministerio Economía y Competitividad, Maria de Maeztu grant (MDM 2014-0431), and Comunidad de Madrid, MadEco-CM (S2015/HUM-3444).

[†]University of Technology Sydney, benjamin.balzer@uts.edu.au

[‡]Carlos III de Madrid, Department of Economics, jschneid@eco.uc3m.es

We show that it is necessary and sufficient to determine the optimal information structure in the fighting stage to determine optimal arbitration.

Our results apply to a large set of environments. We include settings in which information revealed in arbitration has strategic relevance in the fighting stage. Allowing for these cases complicates the problem. Players are aware of the strategic value of information. They have an increased privacy concern and an incentive to manipulate the information structure. Optimal arbitration takes both effects into account. Our results suggest that the information revelation in arbitration is of first-order importance.

In the first part of the paper we develop a conceptual framework for arbitration. We set up a canonical mechanism-design problem of arbitration. We derive it from 5 axioms. We impose: (i) private information about strength in the fight that can be ordered; (ii) efficient settlement; (iii) unilateral veto rights; (iv) exogenous rules of the fighting stage; and (v) arbitration without a structural deficit.

Our framework nests existing models of arbitration, but generalizes beyond these settings.¹ Conceptually, the arbitrator solves two interdependent problems simultaneously. A mechanism-design problem to set the arbitration protocol and an information-design problem to manage beliefs in the fight.

The problem emerges naturally when designing optimal arbitration and fits well in the existing literature. We use it as the basis to derive an alternative formulation of the problem.

Our alternative formulation aims to overcome a drawback of the mechanism-design approach: classic techniques are often impractical (i) to disentangle the mechanism-design part from the information-design part, (ii) to isolate the arbitrator's trade-off, and (iii) to compute the solution.

In the second part of the paper we construct an equivalent formulation of the problem: the belief-management problem. It is an information-design problem directly defined on the fighting stage. Its solution implies the solution to the arbitration problem. Belief management determines the optimal information structure in the fighting stage. We construct a one-to-one mapping from implementable information structures to the candidates for the optimal mechanism. Optimal belief management is necessary and sufficient for optimal arbitration. Belief management allows us (i) to disentangle the mechanism-design part from the information-design part, (ii) to isolate and interpret the arbitrator's trade-off, (iii) and to reduce the computational complexity.

¹Most models in the literature assume either *no or only non-strategic* effects of information revelation on players' behavior in the fighting stage.

We exploit homogeneity of Bayes' rule and the resulting linearity of the problem net of the induced information structure. This linearity is the key step to establish a mapping from any given information structure to a candidate mechanism. That relation holds for any objective in which the arbitrator's budget constraint binds. The precise formulation of the belief-management problem depends on the arbitrator's particular objective. We formulate it for the objectives common in the literature.

Under these objectives, belief-management problems have a clear economic interpretation. The arbitrator seeks to implement fights with a high level of fundamental discrimination and a low level of inefficiency. Fights fundamentally discriminate more if a *ceteris paribus* change in payoff types has a large effect on expected outcomes. Fights have less inefficiencies if the expected aggregate utility is high. The corresponding objective function of the belief-management problem has an economic intuition. It combines direct measures for both goals.

Finally, we apply our results to a set of examples. We cover much of the existing literature as well as more complex settings in which fights are sensitive to the information structure.

Related Literature. The literature on arbitration is extensive. We follow Brown and Ayres (1994) and assume that the main obstacle to settlement is asymmetric information.

One strand of the literature ignores participation constraints and concentrates on a particular class of bargaining protocols (see Armstrong and Hurley, 2002; Olszewski, 2011, and references therein). Mylovanov and Zapechelnyuk (2013) revisit that literature and provide a mechanism-design perspective to it. Our framework nests (a discretized version of) Mylovanov and Zapechelnyuk (2013), but also covers cases in which participation constraints are relevant.

Another, perhaps closer related, strand of the arbitration literature considers cases in which veto constraints are binding. Both the law and economics literature (Bebchuk, 1984; Schweizer, 1989; Spier, 1994, and the literature following), and the literature on international conflicts (Bester and Wärneryd, 2006; Fey and Ramsay, 2011; Jackson and Morelli, 2011; Hörner, Morelli, and Squintani, 2015, and the literature following) consider arbitration mechanisms. A common feature is that the information obtained during arbitration has no effect on players' decisions once settlement fails. Our framework nests these models. In addition, we allow information to affect continuation play.

Models on arbitration where information revelation and continuation play interact are rare. Our own work on alternative dispute resolution (Balzer and Schneider,

2017) considers an all-pay auction as the alternative to arbitration. In that paper we apply the methods developed here (among others) to fully characterize optimal alternative dispute resolution in the legal system.²

Similar to us, models on common agency by Calzolari and Pavan (2006a,b) and Pavan and Calzolari (2009) emphasize that the design choices within a mechanism affect action choices outside that mechanism. The common theme in the large literature on resale (e.g. Gupta and Lebrun, 1999; Zheng, 2002; Goeree, 2003; Carroll and Segal, 2018) is that information revelation affects the outcome of an auction. The literature on aftermarkets (Lauermann and Virág, 2012; Atakan and Ekmekci, 2014; Zhang, 2014; Dworczak, 2017) follows Calzolari and Pavan (2006a) by looking at the interaction between the design of the mechanism and action choices in the aftermarket. However, almost all of that literature makes detailed assumptions on either the mechanism or the disclosure rule. An exception is Dworczak (2017), the paper closest to ours in that literature.

Conceptually, there are two major difference between arbitration and aftermarkets: First, in arbitration the fighting stage serves as the main screening instrument, while the potential to resell an object is often an obstacle to screening. Second, in arbitration the set of players is identical inside and outside the mechanism. As a consequence, in our model all players learn from arbitration about their opponents' type and interpret the information in light of their own history of play. In contrast, in the aftermarket literature players can only learn about those opponents that participated in the mechanism. This restricts the potential information structure under which the aftermarket is played.

On an abstract level the arbitration problem is a mechanism-design problem with adverse selection and moral hazard a la Myerson (1982). It involves information externality in the sense of Jehiel and Moldovanu (2001). Contrary to their paper the externality also effects the players' behavior after the mechanism in our setting.

Our belief-management representation transforms the mechanism design problem into an information-design problem a la Bergemann and Morris (2016). Recent techniques proposed by Mathevet, Perego, and Taneva (2017), Dworczak and Martini (2018), Galperti and Perego (2018), Kolotilin (2018), and Kolotilin and Zapechelnyuk (2019) can be applied to address the belief-management problem.

We relate to these models in another dimension too. Like us, they determine

²Zheng and Kamranzadeh (2018) consider a model identical to the baseline in Balzer and Schneider (2017), but restrict attention to take-it-or-leave-it settlement offers. This arbitration protocol is outperformed by the optimal mechanism characterized in Balzer and Schneider (2017). Zheng (2018) addresses the question when a full settlement mechanism exists given an all-pay auction as fighting game.

the price of implementing an information structure. However, and different to us, they treat the prior distribution as the designer’s initial, exogenous endowment. In our model, the designer endogenously “produces” that endowment. The relation between the mechanism-design part and the information-design part determines the *cost* of producing the “prior” of the fighting stage.³ The arbitrator has a hybrid task: acquire information and disseminate it.

In the above sense closer is recent work by Georgiadis and Szentes (2018) that considers optimal monitoring rules. While different from our model in many dimensions, they, too, obtain an information-design formulation that adds tractability by separating the contract-design part from the information-design part.

Roadmap. In Section 2 we set up a basic model and state our five axioms. In Section 3 we derive a canonical arbitration problem from these axioms. We connect it to existing approaches and document several modifications. In Section 4 we develop our main result. We construct a belief-management problem equivalent to the arbitration problem. We review examples and modifications from Section 3 under belief management. In Section 5 we address difficulties in and solution approaches to the belief-management problem.

2 Model

There are two ex-ante symmetric, risk-neutral players, A and B . Each player i has a privately known *type* $\theta_i \in \Theta \equiv \{1, 2, \dots, K\}$. Types are independently distributed according to a distribution $p : \Theta \rightarrow [0, 1]$ with $\sum_{\theta=1}^K p(\theta) = 1$.

The final outcome depends on which of the following *disjoint events* realizes: veto, \mathcal{V} , escalation, \mathcal{E} , or settlement \mathcal{Z} . We define outcomes separately.

Veto. In the event \mathcal{V} player i receives (exogenous) payoff $V_i(\theta_i) \in (-\infty, 1]$.⁴

Escalation. In the event \mathcal{E} , players play a non-cooperative game. Its rules are a finite set of action profiles \mathcal{A} and a mapping $(u_A, u_B) : \Theta^2 \times \mathcal{A} \rightarrow (-\infty, 1]^2$. We refer to the triple $(\mathcal{A}, u, \Theta^2)$ as the *escalation game*.

Settlement. Settlement, \mathcal{Z} , is an allocation $(x_A, x_B) \in [0, 1]^2$ with $x_A + x_B \leq 1$. In addition, player i receives a utility transfer $t_i \in \mathbb{R}$.

Payoffs and a First Pair of Axioms. Player i type θ_i ’s ex-post payoff is equal to V_i , u_i , or $x_i + t_i$, depending on whether the final outcome is determined through

³Dworczak (2017) determines a similar object (the “no-communication posterior”) for the case of single-agent mechanisms with aftermarkets.

⁴We show later that this includes cases in which payoffs in \mathcal{V} are determined endogenously.

veto, escalation, or settlement. We impose two axioms on the payoffs.

Axiom 1 (Types are ordered). The functions u_i and V_i are increasing in θ_i .

Axiom 2 (Settlement is ex-post efficient). The joint payoff in the event \mathcal{E} never exceeds 1, that is, for all type and action profiles in $\Theta^2 \times \mathcal{A}$ it holds that $u_A + u_B \leq 1$.

Axiom 1 implies that types are ordered. It allows us to interpret θ_i as a player's *strength in the fight*.

Axiom 2 implies that a settlement solution can replicate any (ex-post) outcome of the fight. Conditional on knowing the outcome of fighting, a settlement solution exists that is (weakly) preferred by both players to the fighting outcome. Axiom 2 only concerns the relationship between \mathcal{E} and \mathcal{Z} . Axiom 2 ignores \mathcal{V} because V_i is an ex-interim, reduced-form representation. Thus, an ex-post notion is not well-defined. Later we impose Assumption 1 (page 8) which implies that \mathcal{V} is undesirable ex-post. In Section 3.4 we explicitly model \mathcal{V} (Modification 1). Using this alternative model Axiom 2 applies analogously to \mathcal{V} .

Jointly Axiom 1 and 2 capture the idea of private information as the main driver of conflict. Strong players need to receive a beneficial allocation x_i to be willing to forgo the fight. A weak player may mimic a strong player to have access to these beneficial allocations but faces higher cost when ending up in the fight. Absent adverse selection there is a settlement solution that both parties ex-post prefer over fighting. We now introduce arbitration to the model.

Arbitration and a Second Set of Axioms. Arbitration is a mechanism that players can use to select among the events \mathcal{V} , \mathcal{E} , and \mathcal{Z} . In addition, arbitration can specify direct utility transfers between players. We assume full commitment on all sides once arbitration is accepted by both players.⁵ We impose three axioms on the set of potential arbitration mechanisms.

Axiom 3 (Unilateral Veto Rights). An arbitration mechanism involves unilateral veto rights at the interim stage, that is, each player can (publicly) veto the arbitration mechanism. A single veto triggers event \mathcal{V} .

Axiom 4 (Exogenous Rules of Conflict). An arbitration mechanism treats \mathcal{A} , u_i , and V_i as exogenously given. Players can select any $a_i \in \mathcal{A}_i$ in event \mathcal{E} .

Axiom 5 (Budget Balance). Arbitration cannot run a *structural deficit*, that is, any mechanism is budget balanced in expectation, $\sum_i \sum_{\theta_i} p(\theta_i) t_i(\theta_i) \leq 0$.

⁵Recent literature (Goltsman et al., 2009; Hörner, Morelli, and Squintani, 2015) has explored the role of commitment in models of dispute resolution. While not the focus of our paper our general framework is flexible enough to address these points.

Axiom 3 implies that each player can veto the proposed arbitration mechanism. In that case both parties move directly to event \mathcal{V} .

Axiom 4 implies that arbitration can induce a fight, but cannot control a player's action choice, nor the implied outcome of a fight. We emphasize that Axiom 4 does not necessarily imply that the designer has no influence on the rules of the fighting stage. The arbitrator's choice among different potential games can easily be incorporated in the formulation of u_i and \mathcal{A} . Instead, Axiom 4 allows also for the possibility that only a very specific game u_i is available to the designer. Most of the literature on arbitration assumes a specific escalation game.

Jointly Axiom 3 and 4 capture the idea that arbitration can only control part of the environment. While arbitration has full control over settlement, it has limited control over the fight.

Finally, Axiom 5 implies that the expected increase in joint surplus through arbitration is bounded by 1. We allow the arbitrator to increase the joint surplus beyond that in some cases, but not on average.

Solution Concept. We focus on perfect Bayesian equilibria using the definition of Fudenberg and Tirole (1988).

3 A Canonical Arbitration Problem

In this section we set up a canonical arbitration problem. Formally, the arbitration problem is a mechanism-design problem. It is constrained by the surrounding environment. The arbitrator proposes a game \mathcal{M} that implements either of the events \mathcal{V} , \mathcal{E} , or \mathcal{Z} . While the arbitrator has full control over the event \mathcal{Z} , she ceases control once events \mathcal{V} and \mathcal{E} occur.

We start by defining what we call the *primitives to the arbitrator's problem*. While primitives to the design problem, these terms are derived from our model primitives. Imposing some mild assumptions, we then state and prove the version of the revelation principle that is applicable to our problem.

3.1 Primitives to the Arbitrator

In a perfect Bayesian equilibrium a player best responds to her current information set at any decision node. In particular, suppose we are at a terminal node of the arbitration mechanism. A player's information set consists of a private part and a public part. Player i 's private part, (σ_i, θ_i) , is her exogenous private information, θ_i , and all further information privately acquired previously, σ_i . The public part,

the information structure \mathcal{B} , contains all elements that are public knowledge at the decision node.

Under the notion of perfect Bayesian equilibrium, player i uses her information set to form beliefs. She computes a probability mass function over the opponent's payoff type, $\beta_i : \Theta \rightarrow [0, 1]$ —her belief about the opponent's type. In addition, the player forms a set of conditional probability mass functions $f_{\theta_{-i}} : A_{-i} \rightarrow [0, 1]$ —her belief about θ_{-i} 's continuation strategies. Given her beliefs, the player picks an action that constitutes a best response to these beliefs. These considerations are relevant in the event of escalation. On the equilibrium path players' (conditional) beliefs are correct. The expected continuation value for the (on-path) event \mathcal{E} is

$$U_i(\sigma_i, \theta_i, \mathcal{B}) := \max_{a_i \in A_i} \sum_{\theta_{-i}} \beta_i(\theta_{-i} | \sigma_i, \mathcal{B}) \sum_{A_{-i}} f_{\theta_{-i}}(a_{-i} | \sigma_i, \mathcal{B}) u_i(\theta_i, \theta_{-i}, a_i, a_{-i}), \quad (\text{U})$$

and the function $\beta_i(\theta_i | \cdot)$, $f_{\theta_{-i}}(a_{-i} | \cdot)$ correspond to the correct (on-path) beliefs about the opponent's type and her associated action choices.⁶

From the arbitrator's perspective, the function U_i is a primitive to her problem. Since the rules of conflict are exogenous, the arbitrator can at most influence information sets, but not the shape of the continuation-utility functions.

Treating the function U_i as a primitive we make one of the following two implicit assumptions. Either (i) the equilibrium selection in the continuation game is exogenous and known to the arbitrator or (ii) the arbitrator can control the equilibrium selection and picks the “best” equilibrium given her objective.

Either of these assumptions implies that we focus on implementability of the optimal mechanism. Approaches in which the arbitrator is unaware of the equilibrium-selection rule are beyond the scope of the paper.

The set $\{\Theta, p, U_i, V_i\}$ constitutes the *primitives to the arbitrator*.

3.2 Basic Assumptions

We make two basic assumptions on the functional forms of the arbitrator's primitives. These assumption ensure that there is room for arbitration.

Assumption 1. $U_i(\check{\sigma}_i, \theta_i, \check{\mathcal{B}}) \geq V_i(\theta_i)$ for $\{\check{\sigma}_A, \check{\sigma}_B, \check{\mathcal{B}}\}$ such that $\beta_i(\theta_{-i} | \check{\sigma}_i, \check{\mathcal{B}}) = p(\theta_{-i})$ and $\sigma_i = \emptyset$ for each player.

Assumption 1 implies that an arbitration mechanism with full participation in equilibrium exists. The arbitrator escalates all cases and players' receive at least

⁶The commonly known \mathcal{B} includes the belief closed subset under which the game is played. Players have (some) common prior and there is common certainty of rationality (see Bergemann and Morris, 2016, for further details).

their veto payoffs. Assumption 1 excludes cases in which parties are *exogenously punished* for participation.

Assumption 1 captures cases in which participation in arbitration is neutral or beneficial even if it fails to settle. Benefits may come through psychological components or institutional features such as penalties for not agreeing to arbitration. Examples can be found both in legal systems through explicit sanctions imposed by the court and in international relations where participation in peace talks often implies a temporary lift of imposed sanctions.

We make a second assumption on the arbitrator's set of potential objectives. We assume that the arbitrator's preferences are (weakly) monotone in cases settled and the players' ex-ante expected payoffs.

Assumption 2. Suppose the following is true for two mechanisms \mathcal{M} and \mathcal{M}' .

1. \mathcal{M}' settles at least as many cases as \mathcal{M} ;
2. the (ex-ante) expected payoff of either player in mechanism \mathcal{M}' is at least as large as her payoff in mechanism \mathcal{M} .

Then the arbitrator (weakly) prefers \mathcal{M}' to \mathcal{M} .

Potential objective functions under Assumption 2 include *minimize the ex-ante probability of conflict* and *maximize the player's ex-ante expected utilities*, the objectives mainly used in the literature. Assumption 1 and 2 together with Axiom 1 to 5 define a canonical arbitration problem.

Definition 1 (Canonical Arbitration Problem). A canonical arbitration problem is a mechanism-design problem under Axiom 1 to 5 and Assumption 1 and 2.

We discuss examples of the canonical arbitration problem at the end of the section. Before that, we discuss and state the version of the revelation principle applicable to our environment.

3.3 A Revelation Principle for Optimal Arbitration

To state the revelation principle we define a direct revelation mechanism. A direct revelation mechanism maps the profile of type reports into four objects: (i) a probability, γ , that the conflict escalates to event \mathcal{E} , (ii) a sharing rule, X , determining the allocation when the settlement is achieved, (iii) a direct utility transfer t , and (iv) an additional signal, random variable Σ , with realization (σ_A, σ_B) .

Definition 2 (Direct Revelation Mechanism). A direct revelation mechanism (DRM) is a mapping

$$\mathcal{M}(\cdot) = (\gamma(\cdot), X(\cdot), t(\cdot), \Sigma(\cdot)) : \Theta^2 \rightarrow [0, 1] \times [0, 1]^2 \times \mathbb{R}^2 \times \Delta(\mathcal{S}), \quad (\mathcal{M})$$

where $\Delta(\mathcal{S})$ is the set of distributions over a signal space. For technical reasons that will become clear later we assume without loss that $\mathcal{S} \supseteq \Theta^2 \times \mathcal{A}$, and \mathcal{S} finite. Again, without loss assume that each realization σ_i consists of a tuple, (s_i, m_i) , that returns the player's (own) type report, m_i , and recommends an action choice, s_i .⁷

The mechanism \mathcal{M} is played as follows: Players report their type privately to the mechanism. With probability $\gamma(\theta_A, \theta_B)$ the event \mathcal{E} follows and signal realization σ_i is privately communicated to player i . With the remaining probability $1 - \gamma(\theta_A, \theta_B)$ the event \mathcal{Z} follows and the allocation (x_A, x_B) is implemented. In addition, players receive the utility transfers t in the event \mathcal{Z} . Parties are risk neutral and it does not matter if transfer t is paid only in one event or in all of them. To simplify the exposition we assume it is paid only in event \mathcal{Z} .

Conditional on participation and truthful reporting by $-i$, the expected payoff in a given DRM depends on the type report m_i and the actual type θ_i . It is

$$\begin{aligned} \Pi_i(m_i; \theta_i) = & \underbrace{\sum_{\theta_{-i}} p(\theta_{-i})(1 - \gamma_i(m_i, \theta_{-i}))(x_i(m_i, \theta_{-i}) + t_i(m_i, \theta_{-i}))}_{=: z_i(m_i) \text{ (settlement value)}} \\ & + \underbrace{\left(\sum_{\theta_{-i}} p(\theta_{-i}) \gamma_i(m_i, \theta_{-i}) \right) \sum_{\sigma_i} Pr(\sigma_i | m_i) U_i(\sigma_i; \theta_i, \mathcal{B})}_{=: y_i(m_i; \theta_i) \text{ (escalation value)}}, \end{aligned} \quad (1)$$

with $Pr(\sigma_i | m_i)$ the (conditional) probability that a signal σ_i realizes. The first part, the settlement value, is the expected payoff from event \mathcal{Z} . It depends only on the type report m_i . The second part, the escalation value, is the expected payoff from event \mathcal{E} . It depends on the type report m_i and the actual type θ_i .

A DRM is *incentive compatible* if

$$\forall m_i, \theta_i \in \Theta : \quad \Pi_i(\theta_i; \theta_i) \geq \Pi_i(m_i; \theta_i).$$

A DRM is *incentive feasible* if it is incentive compatible and all types want to participate, that is,

$$\forall \theta_i \in \Theta : \quad \Pi_i(\theta_i; \theta_i) \geq V_i(\theta_i).$$

A DRM is *implementable* if it is incentive feasible and satisfies the arbitrator's budget constraint, that is,

$$\sum_i \sum_{\theta_i} p(\theta_i) \sum_{\theta_{-i}} p(\theta_{-i})(1 - \gamma_i(\theta_i, \theta_{-i})) t_i(\theta_i, \theta_{-i}) \leq 0.$$

⁷The feedback of m_i simplifies notation at a later point.

Using the formulation in (1) it is possible to simplify the description of a DRM such that the arbitrator directly picks a function determining the settlement share z . We refer to this as the reduced-form DRM.

Definition 3 (Reduced-Form Direct Revelation Mechanism). A reduced-form DRM is a collection of mappings

$$\mathcal{M} = (\gamma, z, \Sigma), \quad (\mathcal{M})$$

where γ and Σ are defined as above and $z = (z_A, z_B)$ with each $z_i : \Theta \rightarrow \mathbb{R}$.

We first state that DRMs and reduced-form DRMs are equivalent by reformulating the arbitrator's budget constraint. We use that result to state the appropriate revelation principle.

Lemma 1. *A reduced-form DRM (γ, z, Σ) is implementable if and only if an incentive feasible DRM (γ, X, t, Σ) exists that satisfies*

$$\sum_i \sum_{\theta_i} p(\theta_i) z_i(\theta_i) \leq 1 - \sum_{(\theta_A, \theta_B)} p(\theta_A) p(\theta_B) \gamma(\theta_A, \theta_B). \quad (\text{BB})$$

Proposition 1 (Revelation Principle). *It is without loss of generality to focus on implementable reduced-form DRMs.*

Our setting is dynamic in that a continuation game outside the arbitrator's control interacts with the arbitrator's choices. While it is not straightforward in such settings to use the revelation principle it applies in our framework.⁸

An immediate corollary to Proposition 1 is reminiscent of the findings of Compte and Jehiel (2009) and Zheng (2018).

Corollary 1. *Full settlement is implementable if and only if $V_A(K) + V_B(K) \leq 1$.*

3.4 Alternative Environments and Examples

We address the “canonical” claim of the arbitration problem by showing it is robust to modifications and nests a wide range of models as special cases.

Alternative Environments. We state three modification to our environment and highlight how they influence the description of the problem.

Modification 1 (Veto Leads to Play of a Game). In our model we assume that a veto implies a (possibly type-dependent) exogenous outside option. Here, instead, a veto implies the play of a game $(\mathcal{A}^\nu, v_i, \Theta^2)$ defined analogously to the game in the escalation event; $(v_a, v_b) : \Theta^2 \times \mathcal{A}^\nu \rightarrow (-\infty, 1]^2$. Likewise, we define $V_i(\theta_i, \mathcal{B}^j)$, where \mathcal{B}^j is the public information structure after a veto by $j \in \{i, -i\}$.⁹

⁸Sugaya and Wolitzky (2018) discuss the revelation principle for multistage games in general.

⁹Private information beyond the distribution of types is irrelevant here.

Proposition 2. *If $V_i(\theta_i, \mathcal{B})$ is convex in \mathcal{B} , the arbitration problem is isomorphic to the canonical arbitration problem.¹⁰*

The modification includes parties playing the escalation game $(\mathcal{A}, u, \Theta^2)$ in case of a veto as a special case. That special case is often assumed in the literature (e.g. Schweizer, 1989; Bester and Wärneryd, 2006; Hörner, Morelli, and Squintani, 2015).

Modification 2 (No Transfers). In our model we assume that the arbitrator can impose direct utility transfers between players subject to her budget constraint. Here, instead, we do not allow for direct transfers, that is, $t \equiv 0$.

Without transfers, a reduced-form mechanism (γ, z, Σ) may not be implementable because there is no X that can implement z given γ . To state a necessary and sufficient condition for implementability we follow Border (2007). We define the following terms. For any $Q \subset \Theta^2$ let $Q_i := \{\theta_i \mid \exists \theta_{-i} : (\theta_i, \theta_{-i}) \in Q\}$ and $\tilde{Q} := \{(\theta_A, \theta_B) \in \Theta^2 \mid \theta_i \notin Q_i \text{ for } i = \{A, B\}\}$. Moreover, let $P(\mathcal{E}) := \sum_{(\theta_i, \theta_j)} p(\theta_i)p(\theta_j)\gamma(\theta_i, \theta_j)$. We have the following general implementation condition. For all $Q \subseteq \Theta^2$

$$\sum_i \sum_{\theta_i \in Q_i} z_i(\theta_i)p(\theta_i) \leq 1 - Pr(\mathcal{E}) - \sum_{(\theta_A, \theta_B) \in \tilde{Q}} (1 - \gamma(\theta_A, \theta_B))p(\theta_A)p(\theta_B). \quad (\text{GI})$$

Equations (GI) mirror the general implementation condition from Border (2007).

Proposition 3. *Suppose the mechanism cannot impose direct utility transfers. Then a DRM exists that implements a given reduced-form DRM (γ, z, Σ) if and only if the reduced-form DRM satisfies (GI).*

In the absence of transfers our model is also isomorphic to a model in which a player's type affects her valuation of the share, x_i .

Proposition 4. *Suppose the value of share x_i is $\varphi(\theta_i)x_i$, $\varphi(\theta_i)$ is increasing, and transfers are not available. The arbitration problem is isomorphic to the canonical arbitration problem.*

If transfers are available and a player's valuation for share x_i is $\varphi(\theta_i)x_i$, then the arbitrator does not need the event \mathcal{E} for screening. Instead, she can screen within \mathcal{Z} . The problem collapses to the standard quasi-linear, independent-private-value case.

Modification 3 (Confidential Arbitration). In our model we assume that the arbitrator can send private signals to the players. In particular, she can privately

¹⁰The convexity assumption on V_i is sufficient, but not necessary. A weaker, yet involved assumption is to concentrate on information structures (p, ρ_{-i}) , s.t. ρ_{-i} is the belief player $-i$ holds about i after a veto. With abuse of notation we can replace the convexity assumption by “ $V_i(\cdot, (p, \cdot))$ is on the convex closure of V_i 's graph with respect to p ”.

communicate with one party about her communication with the other party. In some settings such communications may be prohibited. Any message a third-party releases to player $-i$ about information previously received from player i has to be available to player i as well.

The following corollary to Proposition 1 describes that case.

Corollary 2 (Public Signals). *Suppose that arbitration is confidential. Then it is without loss of generality that all signals are public, that is in any realization (σ_A, σ_B) , $s_A \equiv s_B$.*

Other assumptions we can relax without changing results include the following: players' commitment to obey the arbitrator, symmetry, and type independence. However, the relaxations are notationally inconvenient. Our model is identical to one in which budget balance holds on an ex-post level as players are risk neutral (see the arguments in Börgers and Norman, 2009).

While u_i allows *interdependencies*, we exclude *correlation* between the players' type distributions. Correlation would—provided a full-rank condition—free the arbitrator entirely from incentive constraints precisely as in Crémer and McLean (1988). For the case without transfers, correlations would lead to binding (GI) constraints at an otherwise unchanged problem.

Examples. We provide examples of environments captured by our setting.

Example 1 (Exogenous Cost of Conflict). Our first example covers classic settings such as (a discretized version of) the bilateral trade environment of Myerson and Satterthwaite (1983).¹¹ If settlement fails, parties consume an exogenously given outside option which is type-dependent. Thus, u_i is constant in all arguments but θ_i . Consequently, also U_i is constant in all arguments except θ_i . In turn, the beliefs β_i and $f_{\theta_{-i}}$ are irrelevant. In line with Myerson and Satterthwaite (1983) and Compte and Jehiel (2009), combining Axiom 2, Assumption 1, and Corollary 1 implies that full settlement is guaranteed. The result is intuitive: settlement is ex-post efficient.

Example 2 (Conflict as Type-Dependent Lottery). Our second example covers a superset of Example 1. Conflict is a type-dependent lottery, capturing all environments where players' continuation strategies are invariant in the belief $\beta_i(\theta_{-i})$. A player's best response depends on her own type only.¹² The modeling approach is equivalent to e.g. Jehiel and Moldovanu (2001) and Compte and Jehiel (2009). Arbitration problems of that kind can be found in international relations (e.g. Bester and Wärneryd, 2006; Hörner, Morelli, and Squintani, 2015; Meirowitz et al., 2017) and settlement bargaining (e.g. Bebchuk, 1984; Schweizer, 1989).

¹¹Compte and Jehiel (2009), section IIB., provides the necessary relabeling of terms.

¹²The literature refers to these as games with ex-post equilibria (Crémer and McLean, 1985).

Formally, type-dependent lotteries imply that $f_{\theta_{-i}}(a_{-i}|\sigma, \mathcal{B}) = f_{\theta_{-i}}(a_{-i})$ and that equilibrium action $a_i^*(\theta_{-i}, m_i, \sigma_i, \mathcal{B})$ is invariant in (σ_i, \mathcal{B}) . Abusing notation we can re-write

$$U_i(\sigma_i, \theta_i, \mathcal{B}) = \sum_{\theta_{-i}} \beta_i(\theta_{-i}|\sigma_i, \mathcal{B}) u_i(\theta_i, \theta_{-i}).$$

Example 3 (Private Cost of Conflict). The third example is a superset of Example 1 too. Yet, it extends Example 1 along a different dimension by only assuming that u_i is constant in θ_{-i} . That setting covers all cases in which the ex-post payoff of player i depends on the *action profile* and thus on the opponent's *choices*, but not on the opponent's type. The literature on this type of arbitration is small. In a related paper (Balzer and Schneider, 2017) we analyze optimal arbitration in such an environment.¹³ Moreover, the resale literature (e.g. Gupta and Lebrun, 1999; Calzolari and Pavan, 2006a) considers such an environment.

Formally, let $f(a_{-i}|\sigma, \mathcal{B}) := \sum_{\theta_{-i}} \beta_i(\theta_{-i}|\sigma_i, \mathcal{B}) f_{\theta_{-i}}(a_{-i}|\sigma, \mathcal{B})$. Abusing notation again, we can re-write

$$U_i(\sigma_i, \theta_i, \mathcal{B}) = \max_{a_i \in A_i} \sum_{A_{-i}} u_i(\theta_i, a_i, a_{-i}) f(a_{-i}|\sigma_i, \mathcal{B}).$$

4 Belief Management

In this section derive an information-design problem conditional on the escalation game. Solving that problem is necessary and sufficient to find optimal arbitration if full settlement cannot be guaranteed. The necessary part follows from Axiom 4. We show that finding the optimal information structure is sufficient to determine the optimal mechanism too. We focus on cases in which a mechanism implementing full settlement does not exist.¹⁴

Assumption 3 (No full settlement). $V_A(K) + V_B(K) > 1$.

We derive the result in steps. First, we establish our notion of an information structure. Second, we characterize a mapping from implementable information structures to a unique candidate for the optimal mechanism. Third, we apply this result to a range of objectives of the arbitrator. We obtain the corresponding information-design problem. Finally, we revisit the examples from Section 3.4.

¹³Zheng and Kamranzadeh (2018) use a bargaining approach in an otherwise identical model.

¹⁴Assumption 3 is useful to rule out the trivial full-settlement solution when considering a symmetric (i.e., anonymous) objective. For the sake of simplicity, we build our argument ignoring cases in which full settlement is implementable but not optimal. However, nothing changes qualitatively when including those cases expect that the last step of the proof of Theorem 1 becomes slightly more involved.

Our results simplify the arbitration problem. Obtaining the optimal information structure in \mathcal{E} is sufficient to determine the optimal mechanism. That is, for a given canonical arbitration problem we can disentangle the mechanism-design part of eliciting information from the information-design part of distributing that information. The transformation reduces complexity and improves clarity of the arbitrator's (economic) problem. The information-design problem is economically intuitive, highlights the arbitrator's main trade-off, and aides tractability.

4.1 Preliminaries: Information Sets and Notation

The information set at the beginning of the escalation stage consists of a private part and a public part. The private part, (σ_i, θ_i) , contains the realization of the signal, a player's own private history of play, and the exogenous private type. The public part contains everything commonly known, in particular, the arbitrator's choice \mathcal{M} , the information that the conflict escalated, and the common prior.

A player uses all publicly available information to compute her own beliefs about the opponent's potential beliefs and higher-order beliefs. On the equilibrium path these computations are correct. In addition, a player uses her privately obtained information to determine the likelihood that her opponent holds a specific belief (and higher-order beliefs).

Recall that any realization (σ_A, σ_B) provides player i with a private signal $\sigma_i = (s_i, m_i)$. In equilibrium, player i forms a belief about the message σ_{-i} which her opponent receives. That belief itself can be decomposed into two parts. A belief about m_{-i} and one about s_{-i} . We address them in turns.

The belief about m_{-i} coincides with the belief about the payoff type θ_{-i} because \mathcal{M} is incentive compatible. We represent the players' updating procedure as a sequence of two updates. First, the player uses her knowledge of the arbitrator's choice of γ and the observation of event \mathcal{E} . Applying Bayes' rule using the prior p leads to a probability mass function $\beta_i(\theta_{-i}|m_i)$. Second, she uses the realization s_i that may contain additional information about θ_{-i} . Applying Bayes' rule again—using $\beta_i(\theta_{-i}|m_i)$ as the prior—leads to a (refined) belief $\beta(\theta_{-i}|m_i, s_i)$.

We proceed with the belief about s_{-i} . Conditional on report profile (m_A, m_B) , the belief s_{-i} is entirely determined through the known structure of the signal Σ . We represent it by a cumulative distribution function $S_i(s_{-i}|m_i, s_i, \theta_{-i}) : A_{-i} \rightarrow [0, 1]$.

We determine the public information structure *before* the realization of the signal, but *after* the event \mathcal{E} is announced. In that case, each player has sent her type report but has not yet obtained σ_i . The player therefore knows her belief about the opponent's payoff type $\beta_i(\theta_{-i}|m_i)$, a (conditional) distribution over realizations of

her own signal $S_i(s_i|m_i, \theta_{-i}) : A_i \rightarrow [0, 1]$, and a set of (conditional) distributions over her opponent's signal $S_i(s_{-i}|m_i, s_i, \theta_{-i})$. We abuse notation to shorten the description of the latter two to $S_i(\theta_{-i}|m_i) := (S_i(s_{-i}|m_i, s_i, \theta_{-i}), S_i(s_i|m_i, \theta_{-i}))$.

The public information structure contains $(\beta_i(\theta_{-i}|\theta_i), S_i(\theta_{-i}|\theta_i))$ for all types and players. The matrix B with element $b_{ij} := \beta_i(\theta_{-i}|\theta_i)$ describes the belief about the state, (θ_A, θ_B) , up to second order. The matrix S with element $s_{ij} := S_i(\theta_j|\theta_i)$ describes the belief about (expected) realizations of the private signals up to second order. The public information structure at this point is the tuple $\mathcal{B} := (B, S)$.

We determine a set of properties of \mathcal{B} .

Definition 4. A mechanism \mathcal{M} *induces* information structure \mathcal{B} if \mathcal{B} is the public information structure in some on-path continuation game of event \mathcal{E} .

Definition 5 (Consistency). A set of beliefs over the type space, B , is *consistent* with respect to the prior, p , if there is a mechanism \mathcal{M} and a set of beliefs about realizations S such that \mathcal{M} induces (B, S) .

The next two results further qualify consistency.

Lemma 2. B is consistent if and only if there is a γ such that B follows from applying Bayes' rule under γ .

Lemma 3. Take a set of arbitrary probability mass functions each with full support over Θ , $\{\beta_A(\cdot|j)\}_{j \in \Theta} \cup \beta_B(\cdot|1)$. There is a unique set $\{\beta_B(\cdot|j)\}_{j \in \Theta \setminus \{1\}}$ such $B = \{\beta_i(\cdot|j)\}_{j \in \Theta, i \in \{A, B\}}$ is consistent.

We combine the two lemmas above to state a result that allows us to connect our notion of a public information structure to the information-design literature.

Lemma 4. Any information structure can be represented by a tuple (B, Σ) . Moreover, information structure $\mathcal{B} = (B, \Sigma)$ is consistent if and only if B is consistent.

The last result can be interpreted as follows. The continuation game in the event \mathcal{E} is an *incomplete information game* in the sense of Bergemann and Morris (2016). Using their terminology, a basic game consist of four elements. An action set, \mathcal{A} , a payoff function, u , a type space, Θ^2 , and a common prior, B . An incomplete information game is a basic game augmented by a random variable, Σ , determining further, privately held, information.

The first-order and second-order beliefs about the state that players hold when observing the mechanism's outcome are entirely determined by the escalation rule γ as the next corollary to Lemma 2 and 4 shows.

Corollary 3. A reduced-form DRM, (z, γ, Σ) , induces information structure, (B, Σ) if and only if B follows from applying Bayes' rule under γ .

Public Signals. A private signal, Σ , may contain a public component. We have seen above that the realization σ_i can contain a public component that reveals more information about the state and the distribution of signals to all players and types simultaneously. For a given information structure, (B, Σ) , the public component induces a spread over B . A corollary to Lemma 4 states that any realization of the signal can be directly induced by some mechanism.

Corollary 4. *Take any distribution over a set of consistent information structures, $F(\mathcal{B})$. There exists a consistent information structure $\bar{\mathcal{B}} = (\bar{B}, \bar{\Sigma})$ that induces this distribution.*

The special case of public signals implies that Σ contains *only a public component*. An example is confidential arbitration. In that case it is without loss to describe a signal realization by some s , and the players' private signal $\sigma_i = (m_i, s)$. The public information structure after s is $B(s)$. Moreover, Σ describes a mapping from Θ^2 onto a lottery over the set of potential realizations $\{s\}$. Finally, the (expected) common knowledge distribution over types conditional on \mathcal{E} is \bar{B} . Each element is $\sum_{k \in \{s\}} Pr(k|\theta_i) \beta_i(\theta_{-i}|\theta_i, k)$. Any $B(s)$ is consistent, and so is \bar{B} .

4.2 Belief Management

In the previous part we have shown that γ provides the “prior” B to the incomplete information game $(\mathcal{A}, u, \Theta^2, B, \Sigma)$. In this part we show that (B, Σ) is sufficient to determine a unique candidate mechanism (z, γ, Σ) in a canonical arbitration problem. We construct a mapping $M_\Sigma(B) \mapsto (z, \gamma)$. It determines the arbitrator's least-costly option to induce (B, Σ) by a reduced-form DRM.

Definition 6. An information structures is *implementable* if an implementable reduced-form DRM induces it.

Theorem 1. *Consider a canonical arbitration problem. The set of implementable information structures (B, Σ) is compact. Moreover, for any implementable (B, Σ) the optimal reduced-form mechanism, (z, γ, Σ) , is unique.*

Theorem 1 is our main result. The proof is constructive. We organize our discussion of the intuition using a set of observations that correspond to the steps in the formal proof. Recall that B is formed immediately after event \mathcal{E} realized but *before* the signal realizations (σ_A, σ_B) have been communicated.

Observation 1. Both the belief about the state, B , and the continuation payoffs, U_i , are homogeneous of degree 0 in the escalation rule. The escalation value, y_i , is homogeneous of degree 1 in the escalation rule.

Suppose player i submits a report m_i and learns about the event \mathcal{E} . Before receiving any additional signal, the probability of facing a particular type $\tilde{\theta}_{-i}$ is

$$\beta_i(\tilde{\theta}_{-i}|m_i) = \frac{p(\tilde{\theta}_{-i})\gamma(m_i, \tilde{\theta}_{-i})}{\sum_{\theta_{-i}} p(\theta_{-i})\gamma(m_i, \theta_{-i})}.$$

That probability is determined by the *relative* likelihood of escalation only. Thus, if γ implies B so does $\alpha\gamma$. The effect of γ on the continuation payoff in \mathcal{E} is entirely expressed via B . Thus, any U_i is invariant to any scaling of γ . Finally, the (interim) probability of reaching escalation and hence the escalation value are linear in γ , realizations σ are constant in γ .

Observation 2. The *most-costly escalation rule* inducing B is unique.

Take any γ that induces B and pick the largest scalar $\bar{\alpha}$ such that $\bar{\alpha}\gamma(\theta_A, \theta_B) \leq 1$ for all (θ_A, θ_B) . Then, the rule $\bar{\gamma}_B := \bar{\alpha}\gamma$ minimizes the right-hand side of the budget constraint (BB). It is the most-costly rule for the arbitrator. Identifying the most-costly escalation rule is sufficient to characterize *all escalation rules* that induce B . The set of all γ inducing B is $\{\alpha\bar{\gamma}_B : \alpha \in (0, 1]\}$. Given B , the problem reduces to finding the *lowest* α such that (B, Σ) is implementable.

Our next observation contains the main step towards the result.

Observation 3. It is without loss to assume that any type θ_i has either a binding incentive constraint or a binding participation constraint. Given (B, Σ) , all constraints are linear in α and z .

Assumption 2 and 3 ensure that some constraint binds for any type. Otherwise a Pareto improving mechanism with less escalation exists. The second part follows by combining players' expected payoffs with Observation 2.

Observation 3 implies that (B, Σ) captures the entire non-linear part of the constraints. Given (B, Σ) each type has some binding constraint and the set of constraints consists of $2K$ independent linear equations. We have $2K + 1$ unknowns, the $2K$ settlement values and the scalar α . To close the problem we need one more equation. We use the arbitrator's budget constraint (BB).

Observation 4. It is without loss to assume that (BB) holds with equality.

By Observation 3, z_i is linear in α , and thus $\sum_i \sum_{\theta_i} p(\theta_i) z_i(\theta_i)$ is linear too. The right-hand side is independent of z and linear in α . Solving for α delivers a unique tuple $(z_\Sigma^*, \alpha_\Sigma^* \bar{\gamma}_B)$ satisfying the binding constraints with equality for any (B, Σ) .

Via Observation 1 and 2, an implementable escalation rule inducing B exists if and only if the corresponding $\alpha_\Sigma^* \leq 1$.

We construct a function $M_\Sigma : B \mapsto (z, \gamma)$ that (given Σ) identifies a *unique candidate* (z, γ) for any implementable B that can be induced by an implementable \mathcal{M} . It points to the origin otherwise. The function is given by

$$M_\Sigma(B) := \mathbb{1}_{\alpha_\Sigma^* \leq 1} \left(z_\Sigma^*, \alpha_\Sigma^* \bar{g}_B \right),$$

and continuous in the interior of its support.

Discussion of Theorem 1. The canonical arbitration problem contains a mechanism-design part and an information-design part. The arbitrator has full control over the settlement environment, yet under escalation her power ceases and players are free to take decisions.

The informational content of sending players to escalation is endogenous. Moreover, the arbitrator has to evaluate the effects of information revelation on escalation. Thus, it is necessary that the arbitrator solves an information-design problem for that event.

Theorem 1 implies that solving that information-design problem is sufficient too. Given a “prior”, B , and a signal structure, Σ , the optimal mechanism is pinned down by a set of linear equations.

We want to highlight that the arbitrator has more power in the information-design problem implied by Theorem 1 than in Bergemann and Morris (2016). Under Theorem 1, the arbitrator can produce the “prior” to her information-design problem at a cost. The literature on information design assumes that the prior is exogenously given. The reason for that difference is precisely that the arbitrator can control the strategic environment prior to escalation.

The main implication of Theorem 1 is that we can formulate the entire problem focusing only on the event \mathcal{E} . That is, all constraints and the objective are a function of the information structure (B, Σ) only. The set of functions M_Σ determines if such an information structure is implementable and how it is best implemented.

We want to emphasize that the (re-)formulation simplifies the analysis in several ways. Keep in mind that for given “prior”, B , the information-design part of finding the optimal Σ is necessary for any formulation of the arbitration problem. We leave it unchanged in our formulation.

Before this stage our formulation simplifies the problem. On the level of finding a solution, Theorem 1 not only characterizes the set of *implementable* B that some implementable \mathcal{M} can induce, but also limits it to a single candidate mechanism for any B .

Our simplification is independent of the objective itself. The only requirement

we impose is the ‘quasi-Pareto’ criterion defined in Assumption 2. Theorem 1 implies that B is a sufficient statistics for a candidate mechanism for any objective satisfying Assumption 2. To guarantee feasibility we construct a function $M_\Sigma(B)$ that is continuous on the interior of its support. It provides an intuitive mapping from implementable B s to mechanisms, driven by the expected settlement shares necessary to support B .

Focusing on the event \mathcal{E} simplifies the interpretation of both the problem and its results. Using Theorem 1 we can link the properties of the optimal mechanism directly to the information it induces. That is, we can interpret how optimal arbitration manages the information flow between players. The information-design approach allows us to make predictions on the information arbitration reveals.

Our next step is to specify the information-design problem for a class of objective functions of the arbitrator. We argue that the reformulation provides an intuitive description of the main trade-offs and delivers insights even without specifying the escalation game. We then revisit the examples and modifications from Section 3.

4.3 Optimal Belief Management

In this part we apply Theorem 1 to a class of objective functions of the arbitrator. Our formulation covers a range of social welfare functions. It includes all cases in which social welfare is players’ aggregate utility minus some additional cost that escalating the conflict imposes on society. Our result states a well-defined problem of selecting an information structure \mathcal{B} .

We refer to the problem of finding \mathcal{B} as the *belief-management* problem. It is an economically intuitive, compact representation of the arbitration problem, and entirely based on the escalation game $(\mathcal{A}, u, \Theta^2)$.

We start by stating a class of objective functions. Take any $\xi \in [0, 1]$, and assume the arbitrator chooses an implementable mechanism (γ, z, Σ) that solves

$$\max_{(\gamma, z, \Sigma)} (1 - \xi) \underbrace{\left(\sum_i \mathbb{E}[z_i(\theta)] + \left(\sum_{\theta_{-i}} p(\theta_{-i}) \gamma_i(\theta, \theta_{-i}) \right) \widehat{U}_i(\theta; \theta, \mathcal{B}) \right)}_{\text{expected joint surplus}} - \underbrace{\xi \Pr(\mathcal{E})}_{\text{prob. of } \mathcal{E}}, \quad (P_{\mathcal{M}})$$

with

$$\widehat{U}(m_i; \theta_i, \mathcal{B}) := \sum_{\sigma_i} \Pr(\sigma_i | m_i) U_i(\sigma_i; \theta_i, \mathcal{B}).$$

Problem $(P_{\mathcal{M}})$ covers any convex combination between conflict minimization ($\xi = 1$) and joint surplus maximization ($\xi = 0$).

Constraints. A mechanism is implementable if it satisfies participation constraints, incentive constraints and is budget balanced. Theorem 1 states that instead of an implementable mechanism we can directly choose an information structure, $\mathcal{B} = (B, \Sigma)$, such that B is consistent and in the support of $M_\Sigma(B)$. What is left is to determine how to find the *optimal* \mathcal{B} .

To facilitate intuition, think of Σ as (optimally) chosen given prior B . If Σ is optimal conditional on B , the arbitrator's problem is to implement B . However, B not only influences the choice of Σ , but also determines the marginal type distributions conditional on \mathcal{E} , denoted by $(\rho_A(\cdot), \rho_B(\cdot))$. They are the solution to the following system of equations completely identified by B ¹⁵

$$\rho_i(\theta_i) = \sum_{\theta_{-i}} \beta_{-i}(\theta_i | \theta_{-i}) \rho_{-i}(\theta_{-i}) \quad \forall \theta_i, \theta_{-i}. \quad (2)$$

Theorem 1 incorporates the binding participation constraints and budget balance in $M_\Sigma(B)$. What remains is to state the incentive constraints in light of Theorem 1. Take any i and $\theta, \theta' \in \Theta$. Incentive compatibility holds if and only if¹⁶

$$\frac{\rho_i(\theta')}{p(\theta')} \left(\hat{U}_i(\theta'; \theta', \mathcal{B}) - \hat{U}_i(\theta'; \theta, \mathcal{B}) \right) - \frac{\rho(\theta)}{p(\theta)} \left(\hat{U}_i(\theta; \theta', \mathcal{B}) - \hat{U}_i(\theta; \theta, \mathcal{B}) \right) \geq 0. \quad (\text{IC})$$

Objective. To facilitate the exposition we impose more structure on the problem. We impose that apart from the strongest type all types have a (pre-arbitration) incentive to seek settlement.

Assumption 4 (Close Conflicts). $\sum_i \sum_{\theta_i \in \hat{Q}} p(\theta_i) V_i(\theta_i) < \sum_{\theta_i \in \hat{Q}} p(\theta_i)$, for any $\hat{Q} \subseteq \Theta$ and $\hat{Q} \neq \{K\}$.

Assumption 4 imposes structure on the set of relevant constraints.¹⁷

Lemma 5. *Under Assumption 1 to 4 the strongest type's participation constraint holds with equality. All other participation constraints are redundant.*

Jointly, Observation 3 and Lemma 5 imply that we can identify a path, $\iota_i : \Theta \setminus K \rightarrow \Theta$, such that $\Pi_i(\iota_i(\theta_i), \theta_i) = \Pi_i(\theta_i; \theta_i)$, and $\iota_i(\theta) \neq \theta$. The function ι_i determines a binding incentive constraint for each type but the strongest for whom the participation constraint binds by Assumption 3. For the special case that local constraints are sufficient for incentive compatibility $\iota_i(\theta_i) = \theta_i + 1$. Given $\iota_i(\theta_i)$ we

¹⁵The proof of Lemma 3 provides the relevant arguments.

¹⁶For details, see proof of Theorem 1 in particular step 3 and the Lagrangian in appendix B.

¹⁷We derive a general version absent Assumption 4 in the appendix to the paper. The main difference is that it complicates identifying the set of binding constraint.

can define $D_i^\iota(m_i; \theta_i, \mathcal{B}) := \hat{U}_i(m_i; \iota_i(\theta_i), \mathcal{B}) - \hat{U}_i(m_i; \theta_i, \mathcal{B})$. We refer to $D_i^\iota(m_i, \theta_i, \mathcal{B})$ as θ_i 's *type disadvantage* as it describes θ_i 's loss in utility compared to the “next best” type from θ_i 's perspective.

We can use D_i^ι to determine a type's *virtual loss*. Let $\Theta^\iota(\theta) := \{k \in \Theta \mid \exists n \geq 0 \text{ s.t. } \iota^{\circ n}(k) = \theta\}$ be the set of types on path ι that lead to θ after some iteration $\iota^{\circ n}$. Define

$$w_i^\iota(\theta_i) := \sum_{k \in \Theta^\iota(\theta)} \frac{p(k)}{p(\iota(\theta_i))}.$$

For the special case $\iota(\theta) = \theta + 1$ ω^ι describes the hazard rate.

Definition 7 (Virtual Loss). Player θ_i 's virtual loss of pretending

$$\hat{\Psi}_i^\iota(\theta_i, \mathcal{B}) := \begin{cases} w_i^\iota(\theta_i) D_i^\iota(\iota_i(\theta_i); \theta_i, \mathcal{B}) & \text{if } \theta_i \neq K \\ 0 & \text{otherwise.} \end{cases}$$

We state the objective as a function of two objects, each defined within \mathcal{E} .

$$\begin{aligned} \mathbb{E}[\hat{\Psi}|\mathcal{B}] &:= \sum_i \sum_{\theta_i=1}^{K-1} \rho_i(\iota_i(\theta_i)) \hat{\Psi}_i(\theta_i, \mathcal{B}) \quad (\text{expected virtual loss}) \\ \mathbb{E}[\hat{U}|\mathcal{B}] &:= \sum_i \sum_{\theta_i=1}^K \rho_i(\theta_i) \hat{U}_i(\theta_i; \theta_i, \mathcal{B}) \quad (\text{expected utility}) \end{aligned} \tag{3}$$

In principle, multiple incentive constraints could bind. Let $(IC)^A$ be the set of all incentive constraints as defined in equation (IC) which are not governed by ι_i . Define the optimization problem¹⁸

$$\min_B \frac{\xi + (1 - \xi) (1 - \mathbb{E}[\hat{U}|\mathcal{B}])}{\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}] - 1}, \text{ s.t. } (IC)^A \text{ and } M_\Sigma(B) \neq 0. \tag{P_B}$$

Proposition 5 (Duality). *Take a canonical arbitration problem under Assumption 3 and 4. A mechanism solves (P_M) if and only if $(\gamma, z) = M_{\Sigma^*}(B^*)$ and (B^*, Σ^*) solves (P_B) .*

Proposition 5 follows from using the function $M_\Sigma(B)$ to replace z and γ in the arbitrator's objective and rearranging terms. Recall that a consistent B follows from any arbitrary $\{\beta_A(\cdot|j)\}_{j \in \Theta} \cup \beta_B(\cdot|1)$ under Lemma 3.

Proposition 5 provides an intuitive belief-management problem which is equivalent to the canonical arbitration problem. Before discussing the general formulation it is useful to consider the two polar cases. First, if the arbitrator maximizes joint

¹⁸Note that $\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}] > 1$ for any information structure with $M_\Sigma(B) \neq 0$ by Theorem 1.

surplus ($\xi = 0$), she seeks to *maximize*

$$\frac{\mathbb{E}[\hat{\Psi}|\mathcal{B}]}{1 - \mathbb{E}[\hat{U}|\mathcal{B}]}.$$

Second, if the arbitrator minimizes escalation ($\xi = 1$), she seeks to *maximize*

$$\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}].$$

The case $\xi = 1$ has an analogue in revenue-maximizing auction design (Myerson, 1981). The main difference is that—although types are ordered—the term $\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}]$ is non-linear in the arbitrator’s choice. These non-linearities increase complexity. However, conceptually an arbitrator minimizes the likelihood of the conflict by maximizing the expected virtual valuation of the escalation game over the information structure.

More generally, the arbitrator wants to decrease the numerator of the objective in (P_B) . That term captures the cost of escalation to the arbitrator. They consist of a fixed component, ξ , and a variable component, $(1 - \xi)\mathbb{E}(1 - [\hat{U}|\mathcal{B}])$, that captures the (joint) surplus loss of escalation. The cost of escalation are, however, only conditional on the event \mathcal{E} . Therefore, we have to multiply the cost by $(\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}] - 1)^{-1}$. The lower the likelihood of escalation, the lower the expected cost. Thus, the arbitrator wants to maximize the denominator to ensure settlement as often as possible.

We now address the two terms separately. The intuition for the numerator’s form is straight-forward. The higher the welfare losses from escalation, and the more intense the arbitrator’s preferences about it, the more costly is escalation.

The intuition for the denominator’s form is more subtle. To minimize the likelihood of escalation, the arbitrator increases the (expected) virtual loss and utility for players in the event \mathcal{E} . Consider a player in event \mathcal{E} . Increasing her virtual loss contributes to satisfying incentive constraints. The higher $\mathbb{E}[\hat{\Psi}|\mathcal{B}]$, the lower the information rent the arbitrator has to pay to ensure incentive compatibility. Increasing players’ utility, in turn, incentivizes players to agree to participate in arbitration in the first place. It relaxes participation constraints.

Proposition 5 characterizes the economic forces. It provides an intuition how continuation play affects the arbitrator’s choices.

Complexity increases when optimizing the choice of the signal Σ . Nevertheless, combining Proposition 5 and Theorem 1 implies that any remaining complexity is a direct consequence of the associated information-design problem in arbitration.

That is, if we restrict—as most of the literature—complexity such that the scope for information design in the game $(\mathcal{A}, u, \Theta^2)$ is tractable, so is the arbitration problem.

We discuss some avenues of such restrictions in Section 5 and conclude this section by relating the results to the alternative environments and examples discussed in Section 3.4.

4.4 Alternative Environments and Examples

Alternative Environments. We revisit the alternatives from Section 3.

Modification 1 (Veto Leads to Play of a Game). The problems are isomorphic net of the convexity assumption. If convexity is violated, full participation may not be optimal. In our paper Balzer and Schneider (2018) we discuss one way to overcome the participation problem absent the convexity assumption.

Modification 2 (No Transfers). If the arbitrator has no access to additional transfers, she faces the constraints (GI). Net of these constraints the problems are isomorphic by Proposition 3. We recommend the following proceeding when facing a problem in which utility transfers are not possible. When constructing the function $M_\Sigma(B)$ include an additional non-negativity constraint for the settlement values, z .

Unfortunately, and common in the literature on reduced-form mechanism design, it is hard to economically interpret the constraints (GI). We recommend a guess and verify approach. Ignore the constraints (GI) and compute the optimum. If any constraint in (GI) fails, include it and re-optimize.

Modification 3 (Confidential Arbitration). Confidential arbitration reduces the complexity of the signaling space and thus adds structure. It helps us to make statements on the choice, Σ .

Confidential arbitration reduces the set of implementable post-signal information structures to the set of implementable B . As we saw in Section 4.1, Σ reduces to a (potentially type-dependent) lottery over such B 's.

Under confidential arbitration we can apply results from the Bayesian Persuasion literature following Kamenica and Gentzkow (2011). Plain concavification (Aumann and Maschler, 1995) does not apply since lotteries may be type-dependent. However, “Lagrange-Concavification” (Doval and Skreta, 2018) does apply. In Appendix B we state the Lagrangian function of the corresponding constrained *maximization* problem. We show that the optimal Σ is obtained from the Lagrangian’s concave closure.

Finally, there is an intuitive sufficient condition that determines whether signals are needed at all. Fix the signal to be uninformative, that is, $\Sigma(m_A, m_B) :$

$(m_A, m_B) \mapsto ((m_A, \emptyset), (m_B, \emptyset))$. Suppressing the uninformative signal in the notation, we state the following auxiliary problem

$$\min_B \frac{\xi + (1 - \xi)(1 - \mathbb{E}[U|B])}{\mathbb{E}[\Psi|B] + \mathbb{E}[U|B] - 1}, \quad (P_B)$$

with B consistent. Problem (P_B) describes problem (P_B) prohibiting signals and ignoring $(IC)^A$ and $M(B) \neq 0$.

Proposition 6 (No Signals Needed). *Take a canonical confidential arbitration problem under Assumption 3 and 4. If the solution to (P_B) does not violate any constraint $(IC)^A$ and $M(B) \neq 0$ then it is also the solution to (P_B)*

Proposition 6 follows because any consistent B can directly be implemented. If no constraint outside the objective binds, the optimal B is on the convex closure of the objective and the Lagrangian. There is no room for signals to exploit the curvature further.

Examples. We now turn to the examples from Section 3.

Example 1 (Exogenous Cost of Conflict). Recall from the discussion in Section 3 that full settlement is guaranteed in an environment with exogenous cost of conflict through Axiom 2. The belief-management formulation provides another reason for that result. The information channel is entirely shut down. Players consume a given value in case of escalation.

The induced information structure influences neither screening nor welfare in the event \mathcal{E} . Any information structure is thus optimal, but escalation bears some cost. Therefore, full settlement has to be guaranteed. Put differently, Axiom 2 is incompatible with Assumption 3 if u_i is a function of θ_i only.

Example 2 (Conflict as Type-Dependent Lottery). Our second example includes the class of games in which the outcome of \mathcal{E} is determined by a type-dependent lottery. Recall that such environments cover all cases with ex-post equilibria. Different from Example 1, the information structure is relevant. There is an informational externality (Jehiel and Moldovanu, 2001). However, the externality affects continuation utilities linearly because continuation strategies are invariant to the externality. The expected utility U_i is a linear function of B . Thus, constraints and expected values are linear in $\beta_i(\theta_{-i}|\theta_i)/\rho_i(\theta_i) =: \rho(\theta_i, \theta_{-i})$, the joint distribution of type pairs.

Consider escalation minimization ($\xi = 1$). The arbitrator's problem reduces to a linear program of choosing $\rho(\theta_i, \theta_{-i})$. Conditional on the set of binding incentive constraints, any $\rho(\theta_i, \theta_{-i})$ is associated with two constant terms. One captures marginal screening effects, the other captures marginal welfare effects of increasing

the likelihood of a particular profile (θ_A, θ_B) entering \mathcal{E} . Ignoring her budget constraint, the arbitrator puts the entire mass on the the type-profile with the highest combined value. The associated B is degenerate. We can directly verify whether $\rho(\theta_A, \theta_B) = 1$ is implementable by inserting it into a formulation of the budget constraints. If they are not satisfied, we reduce $\rho(\theta_A, \theta_B)$ and the arbitrator assigns the remaining mass in a similar fashion among the remaining type pairs.

Linearity implies symmetry without loss and additional signals never improve upon the no-signal case. In the appendix we formulate an algorithmic solution to the problem under a standard monotonicity assumption. For $\xi < 1$ the general problem becomes equivalent to a quadratic program. It is thus harder to solve.

Yet, most of the literature focuses on escalation minimization. The belief-management approach identifies a condition assumed in most models implying that escalation minimization is equal to joint surplus maximization: The escalation of conflict leads to a constant surplus reduction (see Bester and Wärneryd, 2006; Doornik, 2014; Hörner, Morelli, and Squintani, 2015, among others). That is, the surplus in the event \mathcal{E} is a fixed, type-independent amount $E < 1$. The type-profile determines only *how* E is distributed. In light of Problem (P_B) the consequences are immediate. The objective collapses to $\max \mathbb{E}[\Psi|B]$ and welfare maximization and escalation minimization are equivalent by construction.

Finally, our belief-management formulation provides an alternative take on the solutions of these problems. For example, the binary case of Hörner, Morelli, and Squintani (2015) yields $E[\Psi|B] = (\rho_A(K) + \rho_B(K))\psi$, with constant ψ . The problem reduces to find the highest $\rho_i(K)$ that satisfies the budget constraint.

Example 3 (Private Cost of Conflict). We conclude this section by revisiting our third example. In that example, players' actions are determined by their cost functions *and* by the opponent's expected action. As the belief about the opponent's type changes so does the expected distribution over actions. This, in turn, may trigger a chain of events. Player A adjusts her strategy accordingly, player B responds to that, and so on until a new fixed point is found. Deviations are complex in that setting. Best responses depend on information also off the equilibrium path. In particular, a deviating player can gain herself an information advantage.

To see the information advantage, suppose θ_A deviates by reporting $\theta'_A \neq \theta_A$ and the game moves to event \mathcal{E} . Then θ_A , aware of her deviation, holds belief $\beta_A(\cdot|\theta'_A)$. However, the public information structure remains at its equilibrium value. The opponent, unaware of the deviation, has the same belief as on the equilibrium path. In addition, the opponent's second-order belief is that θ_A holds on-path beliefs as well, and so on. Thus, player B *does not best respond to player A 's action choice*

after player A misreported her type. Instead she best responds to A 's equilibrium action. That reasoning provides an *information advantage* to a deviator. The function Ψ governs how costly (in terms of z) it is to avoid such double deviations. The lower Ψ , the more costly it is.

Equilibrium action choices influence the level of inefficiency in the event \mathcal{E} too. Depending on the information structure, players may choose a costly fight as their continuation path. Alternatively they may choose to concede right away. From the arbitrator's point of view these choices are relevant for two reasons. The discussion of Proposition 5 shows that welfare maximization and minimizing the event \mathcal{E} differ in how much emphasis is put on the efficiency loss, $1 - \mathbb{E}[U|\mathcal{B}]$.

Analyzing Example 3 is considerably harder than Example 1 and 2 and the literature on such problems is sparse. Indeed, analyzing these cases using classic techniques requires to solve a mechanism-design problem with a complex information externality.

Using our results, the problem reduces to an information-design problem. The economic trade-off in that information-design problem is immediate. Although solving that problem cannot be avoided, our results provide guidance towards its solution. In light of our results the mechanism is a simple derivative of that solution.

5 Discussion

Our discussion centers around the issue on how to identify the binding constraints. The formulation in Section 4 provides a problem *conditional on knowing* the binding constraints. The expected utility, U_i , can be a highly non-linear function of the information structure. As a result, it is not easy to identify which constraints bind at the optimum and whether local constraints are sufficient for global constraints.

We provide several results and conditions that help overcoming that tractability issue. We want to emphasize once more that tractability is not a problem of our formulation. To the contrary, our formulation can help to overcome the tractability issues. The difficulty rather stems from the information-design part of the problem. In fact, although information design is conceptually an old question, progress in solving information-design problems on a general level has only been made recently. See, e.g. Mathevet, Perego, and Taneva (2017) and Galperti and Perego (2018) for promising attempts to non-cooperative games.

The difficulty in solving the information-design problem influences the identification of the binding constraints. Thus, at this level of generality we can only provide sufficient conditions. First, we provide a sufficient condition for when local

incentive constraints are sufficient. Let $D_i^+(m, \theta, \mathcal{B})$ be the ability disadvantage if $\iota_i(\theta) = \theta + 1$ for all θ .

Proposition 7. *Local upward incentive constraints imply incentive compatibility if the following holds at the unconstrained optimum*

$$\frac{\rho_i(m_i)}{p(m_i)} D_i^+(m_i; \theta_i, \mathcal{B}) \quad \text{is non-decreasing in } m_i. \quad (4)$$

Even if condition (4) is not satisfied we provide a sufficient condition for local upward incentive compatibility.

Definition 8 (MDR). The game $(\mathcal{A}, u, \Theta^2, \mathcal{B})$ satisfies the monotone difference ratio condition **(MDR)** if $D_i^+(m; \theta_i, \mathcal{B})/D_i^+(m-1; \theta_i, \mathcal{B})$ is non-decreasing in θ_i .

Proposition 8. *Suppose (MDR) holds at the optimum. Local incentive constraints imply (global) upward incentive compatibility.*

Consider the following algorithmic guess and verify approach. First, solve problem $(P_{\mathcal{B}})$ assuming $\iota(\theta) = \theta + 1$. If the solution satisfies (MDR), check if additional downward incentive constraints in $(IC)^A$ are violated. If so, use these constraints to replace one belief in B and solve over the constraint set. Do so until you have found an optimum. Given that (MDR) holds, the algorithm provides a solution.

Independently of whether (MDR) is satisfied, if the optimal solution assuming $\iota(\theta) = \theta + 1$ is monotone in the sense of condition (4), ignored incentive constraints are redundant. While monotone solutions appear intuitive, they cannot be guaranteed for all games. We provide two conditions on the primitives of the escalation game implying monotone solutions. Each condition assumes a *type-separable escalation game*. The payoff function of such a game takes the following form

$$u(a_i; a_{-i}, \theta_i) = \phi(a_i, a_{-i}) - \zeta(\theta_i)c(a_i, a_{-i}),$$

with $\zeta > 0$ decreasing, ϕ, c positive, and strictly increasing in a_i . Moreover, we assume that ϕ is decreasing in a_{-i} and u is concave in a_i . For the sake of the argument, we consider the limiting case of a convex action space. Further we assume that ϕ, c are twice differentiable.¹⁹

Constant Difference Ratio. Condition (4) holds if the ratio $D^+(m_i; \theta_i, \mathcal{B})/D^+(m_i-1; \theta_i, \mathcal{B})$ is constant in θ_i . Incentive compatibility is satisfied if and only if the expected escalation probability $\gamma_i(m_i)$ is non-decreasing in m_i , which follows from a monotone hazard rate given a constant difference ratio.

¹⁹Formally, we use our model and assume the distance between any two actions approaches 0. The limit result allows us to exploit envelope arguments and a representation in compact notation.

Proposition 9. *The difference ratio is constant if, for given distribution of the opponent's action, the (expected) cost function of a player's best response is separable, that is, $\mathbb{E}[c(a_i(m_i; \theta_i, \mathcal{B}), a_{-i}) | m_i, \mathcal{B}] = h(m_i; \mathcal{B})\tilde{g}(\theta_i)$.*

A simple example of such a game is to assume an action space $\mathcal{A} = [0, 1]^2$, $\phi(a_i, a_{-i}) = 1/2 + a_i(1 - a_{-i}) - a_{-i}$, and $c(a_i, a_{-i}) = a_i^2$. Moreover, let $\zeta(\theta) = 1/\theta$. For given distribution of her opponent's action, a player's best response is $a_i(m_i; \theta_i, \mathcal{B}) = (1 - \mathbb{E}[a_{-i} | m_i, \mathcal{B}])\theta_i/2$ which is separable, and so is c .

Non-Constant Difference Ratios. If the difference ratio is non-constant we can specify sufficient conditions. For simplicity we assume that $\xi = 1$ and $\Delta\theta := \zeta(\theta - 1) - \zeta(\theta)$ sufficiently small. The main ingredients to the model to guarantee a monotone solution is that actions are strategic complements. Suppose further that the function ϕ provides a division of the pie, that is, $\phi(a_i, a_{-i}) + \phi(a_{-i}, a_i) = 1$, best responses are continuous, and the hazard rate, $\omega(\theta_i) := \sum_{k=1}^{\theta_i} p(k)/p(\theta_i)$, is non-decreasing. An algorithm close to the one solving Example 2 yields the optimal solution. In appendix D we sketch that algorithm. A simple parameterization is $\phi(a_i, a_{-i}) = 1/2(1 + a_i - a_{-i})$, and $c(a_i, a_{-i}) = (a_i)^2 + 2Ka_i(1 - a_{-i})$.

Appendix

A Proofs

A.1 Proof of Lemma 1

Proof. From the definition of z it follows that (γ, z, Σ) is incentive feasible if and only if (γ, X, t, Σ) is incentive feasible, where (γ, X, t) implies z .

Necessity. The ex-ante expectations of settlement values,

$$\sum_i \sum_{\theta_i} p(\theta_i) \sum_{\theta_{-i}} p(\theta_{-i}) (1 - \gamma_i(\theta_i, \theta_{-i})) x_i(\theta_i, \theta_{-i}),$$

cannot exceed the ex-ante probability of settlement

$$1 - \sum_{(\theta_A, \theta_B)} p(\theta_A) p(\theta_B) \gamma(\theta_A, \theta_B).$$

Implementability implies

$$\sum_i \sum_{\theta_i} p(\theta_i) \sum_{\theta_{-i}} p(\theta_{-i}) (1 - \gamma_i(\theta_i, \theta_{-i})) t_i(\theta_i, \theta_{-i}) \leq 0.$$

Together that implies (BB).

Sufficiency. Suppose (z, γ, Σ) is incentive feasible and satisfies (BB). Let $x_A(\theta_i, \theta_{-i}) = 1$ and pick t_A such that

$$z_A(\theta_A) = \sum_{\theta_B} p(\theta_B) (1 - \gamma_A(\theta_A, \theta_B)) x_A(\theta_A, \theta_B) + t_A(\theta_A).$$

Further, pick $t_B(\theta_B) = z_B(\theta_B)$. Then,

$$\sum_i \sum_{\theta_i} p(\theta_i) \sum_{\theta_{-i}} p(\theta_{-i}) (1 - \gamma_i(\theta_i, \theta_{-i})) t_i(\theta_i, \theta_{-i}) \leq 0$$

and (z, γ, Σ) is implementable. \square

A.2 Proof of Proposition 1

Proof. Take an arbitrary arbitration mechanism. At a terminal node of this game outcome \mathcal{Z} or \mathcal{E} realizes and this information is common knowledge among players. Moreover, if outcome \mathcal{E} realized a player uses her private history of play and the public history to draw inference about her opponent's private type and her private history (e.g., higher order beliefs). Then, players use this information to play \mathcal{E} . That is, the information structure induced by the play of the arbitration mechanism influences players' utility from outcome \mathcal{E} .

From an ex-ante perspective, equilibrium play in the arbitration mechanism implies a type-dependent distribution over outcomes \mathcal{Z} or \mathcal{E} together with a realized private and public history. Each type of a player has access to another types' distribution by imitating her equilibrium strategy. In equilibrium, this is not beneficial. Thus, this distribution can be directly implemented by an incentive feasible DRM. Depending on the report profile, the arbitrator decides which event occurs and sends a potentially random private signal to each player. Players use the signals together with the knowledge about their report, the design of γ , and the function that maps reports into private signals to update the information structure.

Restricting attention to reduced-form mechanisms is without further loss by Lemma 1. Full participation is optimal by Assumption 1. Any on-path veto outcome $V(\theta_A) + V(\theta_B)$ can be replicated inside a reduced-form DRM. \square

A.3 Proof of Corollary 1

Proof. Full settlement implies pooling. A full settlement solution is incentive feasible iff z_i constant and weakly larger than $\max_{\theta \in \Theta} V_i(\theta)$. V_i is decreasing and a solution exists iff $V_A(1) + V_B(1) \leq 1$ by (BB). \square

A.4 Proof of Proposition 2

Proof. Axiom 3 implies that the identity of the vetoing players becomes common knowledge before players play the veto game.

Fix \mathcal{M} and an equilibrium that implies full participation. Then, \mathcal{B} is relevant only off path. In addition, \mathcal{B} is such that the vetoing player holds his prior beliefs about the non-vetoing player by the don't-signal-what-you-don't-know condition of PBE.

On path vetoes are only relevant if they facilitate participation by others. Participation of i is facilitated if $-i$'s vetoes lower i 's expected $V(\theta_i, \mathcal{B})$. By Jensen's inequality that can only happen if V is non-convex in \mathcal{B} (for formal arguments see Celik and Peters, 2011; Balzer and Schneider, 2018, in particular Proposition 2 in the former and Section 3 in the latter). \square

A.5 Proof of Proposition 3

Proof. The proof directly follows from Border, 2007, Theorem 3. \square

A.6 Proof of Proposition 4

Proof. Assume that players value the share according to $\varphi(\theta_i)$. We show that the model is identical to the main model.

An Auxiliary Model. Assume valuations are reversely ordered, that is, $\theta_1 > \theta_2 > \dots > \theta_K$. Define $\check{U}_i(m_i; \theta_i, \mathcal{B}) := U_i(m_i; \theta_i, \mathcal{B})/\varphi(\theta_i)$, where U is the (continuation) payoff from the escalation game. Similarly, transform the outside options $\check{V}_i(\theta_i) := V_i(\theta_i)/\varphi(\theta_i)$ and the values from participating in the mechanism

$$\check{\Pi}_i(m_i; \theta_i) := \Pi(m_i; \theta_i)/\varphi(\theta_i) = z_i(m_i) + \left(\sum_{\theta_{-i}} p(\theta_{-i}) \gamma_i(\theta, \theta_{-i}) \right) \check{U}_i(m_i; \theta_i, \mathcal{B}).$$

Using the transformed terms only the model is identical to that in the main text. We call it the auxiliary model. What remains to show is that the auxiliary model's solution implies the correct model's solution.

Participation Constraints. If the auxiliary model implies full participation, so does the correct model because

$$\check{\Pi}_i(\theta_i; \theta_i) \geq \check{V}_i(\theta_i) \Leftrightarrow \theta_i \cdot \check{\Pi}_i(\theta_i; \theta_i) \geq V_i(\theta_i) \Leftrightarrow \Pi_i(\theta_i; \theta_i) \geq V_i(\theta_i).$$

Full-settlement in the auxiliary model implies

$$\check{\Pi}_i(1; 1) \geq \check{V}_i(1) \Leftrightarrow z_i(1) \geq \check{V}_i(1) \Leftrightarrow \theta_1 \cdot z_i(1) \geq V_i(1).$$

Resource feasibility requires $\sum_i \check{V}_i(1) \leq 1$. Hence full-settlement in the correct model is implementable if and only if it is implementable in the auxiliary model.

Incentive Constraints. Incentive compatibility in the auxiliary model implies incentive compatibility in the correct model because

$$\check{\Pi}_i(\theta_i; \theta_i) - \check{\Pi}_i(m_i; \theta_i) \geq 0 \Leftrightarrow \theta_i \left(\Pi_i(\theta_i; \theta_i) - \Pi_i(m_i; \theta_i) \right) \geq 0 \Leftrightarrow \Pi_i(\theta_i; \theta_i) - \Pi_i(m_i; \theta_i) \geq 0.$$

Budget balance holds because all z are identical in both models. \square

A.7 Proof of Lemma 2

Proof. Take a consistent B . Then, there is a mechanism that implements (B, S) for some S . Thus, there is γ such that B follows from Bayes' rule given γ . Moreover, take any B that follows from Bayes' rule given γ . Then, let $\sigma_i = (\theta_i, \emptyset)$. The mechanism implements (B, S) with $S = 1$. \square

A.8 Proof of Lemma 3

Proof. If all beliefs in $\{\beta_A(\cdot|j)\}_{j \in \Theta} \cup \beta_B(\cdot|1)$ have full support the proof is a direct application of Bayes' rule.

Assuming full support, Bayes' rule implies that

$$\beta_i(\theta_{-i}|\theta_i) = \frac{p(\theta_i)p(\theta_{-i})\gamma(\theta_i, \theta_{-i})}{Pr(\mathcal{E})\rho_i(\theta_i)}, \quad (5)$$

with $\rho_i(\theta_i) := Pr(\theta_i|\mathcal{E})$ the likelihood that player i has type θ_i conditional the event \mathcal{E} . Consistency implies that

$$\rho_i(\theta_i) = \sum_{\theta_{-i}} \beta_{-i}(\theta_i|\theta_{-i})\rho_{-i}(\theta_{-i}). \quad (6)$$

Now fix a set of probability mass functions $\{\beta_A(\cdot|j)\}_{j \in \Theta} \cup \beta_B(\cdot|1)$ each with full

support over the entire type space Θ , but otherwise arbitrary.

Using equation (6) and $\{\beta_A(\theta_B|j)\}_{j \in \Theta}$ we can express any $\rho_B(\theta_B)$ as a (linear) function of the vector $(\rho_A(j))_{j \in \Theta}$. Applying equation (5) to $\beta_B(\theta_A|\theta_B)$ and $\beta_A(\theta_B|\theta_A)$ implies that

$$\rho_B(\theta_B)\beta_B(\theta_A|\theta_B) = \beta_A(\theta_B|\theta_A)\rho_A(\theta_A). \quad (7)$$

Applying equation (7) for $\theta_B = 1$ and any $\theta_A \in \Theta$, and substituting for $\rho_B(1)$ as a function of $(\rho_A(j))_{j \in \Theta}$ determines $(\rho_A(j))_{j \in \Theta}$, and thus $\rho_B(\theta_B)$.

Finally, using equation (5) once more determines the remaining functions $\beta_B(\theta_A|\beta_B \neq 1)$ uniquely.

We want to stress that the proof requires that the initial set of beliefs has full support. Yet, we discuss in the proof of Theorem 1 that restricting ourselves to (limits of) sets with full support is sufficient for the analysis. \square

A.9 Proof of Lemma 4

Proof. Given the explanation in the main text directly below Lemma 4, the proof follows from Bergemann and Morris (2016). \square

A.10 Proof of Theorem 1

We prove Theorem 1 in steps. Steps 1–4 correspond to Observations 1–4 in the main text. Step 5 proves compactness of the choice set. By Corollary 3, it is sufficient to show that for any Σ there is a one-to-one mapping between the reduced-form mechanism and the “prior” B .

Step 1: Homogeneity. We show that B is homogeneous of degree 0 w.r.t. γ via the following claim.

Claim. γ implements B iff every escalation rule $\hat{g}_B = \alpha\gamma$ implements B where α is a scalar.

Proof. Suppose γ implements B . Homogeneity of Bayes’ rule implies that any escalation rule $\hat{g}_B = \alpha\gamma$ implements B . For the reverse suppose $\alpha\gamma$ implements B and set $\alpha = 1$. If γ is an escalation rule it implements B . \square

If B is homogeneous of degree 0 w.r.t. γ so is U_i ; γ is homogeneous of degree 1 by definition and so is y_i .

Step 2: Most-Costly Escalation Rule. We show that B , for fixed Σ , determines $Pr(\mathcal{E})$. That is, the set of all escalation rules implementing a given information

structure, (B, Σ) , is defined up to the real numbers $\{\alpha\}$. The escalation probability is linear in any α .

Fix a consistent B and take *some* escalation rule $\hat{\gamma}$ that implements B . Step 1 implies that each escalation rule that implements B satisfies

$$Pr(\mathcal{E}) = \sum_{(\theta_A, \theta_B)} p(\theta_A)p(\theta_B)\alpha\hat{\gamma}(\theta_A, \theta_B).$$

Let $\{\alpha\}$ be the set of all α such that $\forall(\theta_A, \theta_B)$, $\alpha\hat{\gamma}(\theta_A, \theta_B) \leq 1$ and $\hat{\gamma}(\theta_A, \theta_B) = \alpha\hat{\gamma}(\theta_A, \theta_B) \leq 1$. The set $\{\alpha\}$ determines all escalation rules implementing B . Its largest element determines the most-costly escalation rule uniquely.

Step 3a: Set of binding Constraints and Linearity in $\{\alpha\}$. Consider the optimal mechanism.

Claim. For any θ_i the participation constraint or an incentive constraint is satisfied with equality.

Proof. To the contrary, suppose neither the participation constraint nor an incentive constraint holds with equality. Then, we can reduce $z_i(\theta_i)$ until one of the above constraints holds with equality, and all constraints remain satisfied. \square

Let Θ^{IC} be the set of types with at least one binding incentive constraint and let Θ^{PC} be the set of types with a binding participation constraint. By the previous claim $\Theta^{IC} \cup \Theta^{PC} = \Theta$. Further, let $\Theta^I(\theta_i)$ be the set of types such that θ_i 's incentive constraints w.r.t to any $\theta \in \Theta^I(\theta_i)$ hold with equality. We say $\hat{\Theta}_i \subset \Theta_i^{IC}$ describes a *cycle* if for any $\theta_i \in \hat{\Theta}_i$, it holds that $\theta_i \notin \Theta_i^{PC}$ and $\Theta_i^I(\theta_i) \subset \hat{\Theta}_i$.

Claim. It is without loss of generality to assume no cycles exist.

Proof. Suppose $\hat{\Theta}_i$ describes a cycle. Reducing $z_i(\theta_i)$ for all $\theta_i \in \hat{\Theta}_i$ under condition $z_i(\theta_i) - z_i(\theta'_i) = y_i(\theta'_i; \theta_i) - y_i(\theta_i; \theta_i)$ for any $\theta' \in \Theta^I(\theta_i)$ is possible without violating any other constraint since $\Theta_i^I(\theta_i) \cap \{\Theta_i^{PC} \cup \{\Theta_i^I(k)\}_{k \notin \hat{\Theta}_i}\} = \emptyset$. \square

Claim. z_i is linear in α given B .

Proof. Consider $\theta_i \in \Theta_i^{PC}$. Then, $z_i(\theta_i) = V_i(\theta_i) - y_i(\theta_i; \theta_i)$. The first term of the RHS is a constant, the second is linear in α by step 1. For any $\theta_i \in \Theta_i^{IC}$, the incentive constraint is $z_i(\theta_i) = z_i(\theta'_i) + y_i(\theta'_i; \theta_i) - y_i(\theta_i; \theta_i)$ if $\theta'_i \in \Theta_i^I(\theta_i)$. Given $z_i(\theta'_i)$, linearity holds because y_i is linear in α by step 1. Now, either $\theta'_i \in \Theta_i^{PC}$, or, $z_i(\theta'_i)$ is linear given some $z_i(\theta''_i)$ with $\theta''_i \in \Theta_i^I(\theta_i)$. No cycles exist so that recursively applying the last step yields the desired result. \square

Step 3b: Homogeneity of the expected Shares. Using the results from step 3a, let $\mathbb{P}_i(\Theta)$ describe the finest partition of Θ into subsets Θ_i^p such that for every $\theta_i \in \Theta_i^p \in \mathbb{P}_i(\Theta)$ every $\Theta^I(\theta_i) \subseteq \Theta_i^p$. Let $\Theta_i^p(\theta) := \{\Theta_i^p \in \mathbb{P}_i(\Theta) : \theta \in \Theta_i^p\}$ identify

the element of the partition to which θ belongs. Finally, let $\hat{\theta}_i(k) := \{\max \theta \in \Theta_i^P(k) : \theta \in \Theta_i^P\}$. By the first two claims of step 3a it is without loss to assume that all objects are well-defined and thus $\hat{\theta}_i$ is non-empty for any $\theta_i \in \Theta$. Using the last claim in Step 3a, we can find a set of functions $H_i(\gamma)$ solving

$$\sum_{\theta_i} p(\theta_i) z_i(\theta_i) = -H_i(\gamma) + \sum_{\theta_i \in \Theta_i^{PC}} p(\theta_i) V_i(\theta_i) + \sum_{\theta_i \in \Theta_i^{IC}} p(\theta_i) V_i(\hat{\theta}_i(\theta_i)). \quad (8)$$

Straightforward algebra implies $H_i(\alpha\gamma) = \alpha H_i(\gamma)$. Thus, $H_i(\alpha\gamma)$ is homogeneous of degree 1 in γ .

Step 4: Determining α via constraint (BB) An arbitration outcome is only implementable if the ex-ante expected settlement values are weakly lower than the probability of settlement, (BB). That is, $\sum_i \sum_{\theta_i} p(\theta_i) z_i(\theta_i) \leq 1 - Pr(\mathcal{E})$, where the RHS is strictly lower than 1 by Assumption 3. By step 1 any escalation rule $\alpha\gamma$ implements the same B . If each $\alpha\gamma$ is implementable then $\alpha\gamma$ satisfies (BB). By step 3b we can rewrite (BB) as

$$\sum_{\theta_i \in \Theta_i^{PC}} p(\theta_i) V_i(\theta_i) + \sum_{\theta_i \in \Theta_i^{IC}} p(\theta_i) V_i(\hat{\theta}_i(\theta_i)) - 1 \leq \sum_i \alpha H_i(\gamma) - Pr(\mathcal{E}). \quad (BB')$$

Given Θ_i^{PC} , Θ_i^{IC} , and $\{\hat{\theta}_i(\theta_i)\}_{\theta_i}$ the LHS is independent of the arbitrator's choice. Moreover, the LHS must be positive at the optimum because $\alpha \rightarrow 0$ implies convergence to full settlement which is ruled out by Assumption 3. An immediate result of that is that $\sum_i \alpha H_i(\gamma) > 0$.

If the arbitrator lowers α all constraints continue to hold, but $Pr(\mathcal{E})$ decreases. Redistributing these resources equally among all player types improves upon the proposed mechanism in all three dimensions imposed on the arbitrator's objective by Assumption 2. Thus, such a mechanism cannot be optimal. Equation (BB') and thus equation (BB) hold with equality at the optimal mechanism.

Step 5: Compactness of (B, Σ) .

Any signal, s , can be implemented and since the signal space is finite, the set of signals is closed and bounded. For B , the proof of Lemma 3 provides the main argument. The set of full support distributions $\{\beta_A(\cdot|j)\}_{j \in \Theta} \cup \beta_B(\cdot|1)$ is convex, and so is the set of B s having full support. The next step is to show that the closure of that set can be attained as well. We show this by using the fact that any B on the closure of the set of (full support) B s can be attained by a sequence of full support B s, $\{B_n\}_{n \in \mathbb{N}}$. The following two lemmas provide that result and thus compactness of (B, Σ) . Moreover, all constraints are weak inequality constraints, such that the set of implementable (B, Σ) is compact.

Lemma 6. *Any B is consistent if and only if it can be approximated by a convergent sequence of consistent B with full support.*

Proof. Take a sequence of consistent $B_n \rightarrow B$. B_n is consistent, that is, it implies some function $f : B \rightarrow [0, 1]^{K \times K}$, such that $f(B_n) = \gamma_n$ with γ_n implementing B_n . Since f is continuous, $\lim_{n \rightarrow \infty} f(B_n) = f(\lim_{n \rightarrow \infty} B_n) = \gamma$. Consistency implies

$$g_L(B) = g_R(B), \quad (9)$$

where both g_L and g_R are continuous functions $B \rightarrow \mathbb{R}$. We can conclude that $g_L(B) - g_R(B) = \lim_{n \rightarrow \infty} [g_L(B_n) - g_R(B_n)] = 0$ and B satisfies consistency. This holds because $g_L(B_n) - g_R(B_n) = 0$.

Conversely, take any B that is implemented by some γ . We show that we can find a sequence of interior B that are consistent and converge to B : Let $\hat{\gamma}$ be the escalation rule that implements B . Choose a sequence of escalation rules in the interior that converges to $\hat{\gamma}$. By Bayes' rule every element of the sequence, γ_n , corresponds to some B_n . Moreover, consistency implies that there exists a continuous function, say $f^{-1} : [0, 1]^{K \times K} \rightarrow [0, 1]^{K \times K}$, such that $f^{-1}(\gamma_n) = B_n$ and B_n satisfies consistency. Note that f^{-1} is continuous which implies that $\lim_{n \rightarrow \infty} B_n = \lim_{n \rightarrow \infty} f^{-1}(\gamma_n) = f^{-1}(\hat{\gamma}) = B$. \square

Lemma 7. *Let \mathcal{O} be a continuous function defined on the domain of $\gamma \in [0, 1] \cap C$ where C consists of those γ 's that satisfy a given set of weak inequality constraints, each of which is continuous in γ . Then, $\arg \max_{\gamma \in [0, 1] | \gamma \in C} \mathcal{O}(\gamma) = \arg \sup_{\gamma \in (0, 1) | \gamma \in C} \mathcal{O}(\gamma)$.*

Proof. Without loss of generality suppose the argument that maximizes \mathcal{O} , γ^* , gives rise to a non-interior B , B^* . Then, Lemma 6 implies that we can approximate B^* by a convergent sequence of consistent interior B s. It follows that $\lim_{n \rightarrow \infty} \mathcal{O}(B_n) = \mathcal{O}(B^*)$ because \mathcal{O} is continuous in γ (and through Observation 1 continuous in B). Moreover, because the constraints are inequality constraints and continuous in γ (and B), there is n' such that every element B_n with $n > n'$ satisfies the constraints. Therefore, $\max_{\gamma \in [0, 1] | \gamma \in C} \mathcal{O}(\gamma) = \sup_{\gamma \in (0, 1) | \gamma \in C} \mathcal{O}(\gamma)$ and $B^* = \lim_{n \rightarrow \infty} B_n$. Using Lemma 6 we note that for every B_n there is γ_n so that $\lim_{n \rightarrow \infty} \gamma_n = \gamma^*$. \square

A.11 Proof of Lemma 5

Proof. Consider the resource constraint. Focus on the formulation (BB') in the proof of Theorem 1, step 4. Assume by contradiction that the set of types with binding participation constraint $\Theta_i^{PC} \neq \{K\}$. The LHS of (BB') is $\sum_{\theta_i \in \Theta_i^{PC}} p(\theta_i) V_i(\theta_i) +$

$\sum_{\theta_i \in \Theta_i^{IC}} p(\theta_i) V_i(\hat{\theta}_i(\theta_i)) - 1$. By Assumption 4 this is negative if $\Theta_i^{PC} \neq \{K\}$, contradicting Assumption 3. \square

A.12 Proof of Proposition 5

Proof. All references to *steps* refer to those in the proof of Theorem 1. The rule \bar{g}_B is defined in step 2, H_i and $\hat{\theta}(\theta)$ are defined in step 3. Further define $\gamma_i(m_i) := \sum_{\theta_{-i}} p(\theta_{-i}) \gamma_i(m_i, \theta_{-i})$.

At the optimum (BB') holds with equality (step 4). Substituting for $\sum_{\theta_i} p(\theta_i) z_i(\theta_i) = (1 - Pr(\mathcal{E}))$ in problem (P_M) implies objective

$$(1 - \xi) \left(1 + \sum_i \sum_{\theta} p(\theta) \gamma_i(\theta) \hat{U}_i(\theta; \theta, \mathcal{B}) \right) - \xi Pr(\mathcal{E}). \quad (10)$$

Bayes' rule implies that $p(\theta) \gamma(\theta) = \rho(\theta) Pr(\mathcal{E})$. Factoring out $-Pr(\mathcal{E})$ yields

$$-Pr(\mathcal{E}) \left(\xi - (1 - \xi) \sum_i \sum_{\theta} \rho_i(\theta) \hat{U}_i(\theta; \theta, \mathcal{B}) \right) = -Pr(\mathcal{E}) \left(\xi - (1 - \xi) \mathbb{E}[\hat{U}|\mathcal{B}] \right).$$

For a given B step 1 and 2 imply $\gamma(\theta_A, \theta_B) = \alpha \bar{g}_B(\theta_A, \theta_B) \Rightarrow Pr(\mathcal{E}) = \alpha R(B)$ with $R(B) := \sum_{\theta_A \times \theta_B} p(\theta_A) p(\theta_B) \bar{g}_B(\theta_A, \theta_B)$. (BB') is binding, thus (step 4)

$$\alpha = \frac{Pr(\mathcal{E}) + C}{\sum_i H_i(\gamma)} \quad (11)$$

with constant $C = \sum_{\theta_i \in \Theta_i^{PC}} p(\theta_i) V_i(\theta_i) + \sum_{\theta_i \in \Theta_i^{IC}} p(\theta_i) V_i(\hat{\theta}_i(\theta_i)) - 1$. Substituting for α in $Pr(\mathcal{E}) = \alpha R(B)$ using equation (11) and rearranging implies

$$1 + \frac{C}{Pr(\mathcal{E})} = \frac{\sum_i H_i(\gamma)}{R(B)},$$

and the right hand side is (by step 3 and 4)

$$\sum_i H_i(\gamma) / R(B) = \sum_i \left(\sum_{\theta=1}^{K-1} \rho_i(\iota_i(\theta)) w_i^t(\theta_i) D_i^t(\iota_i(\theta); \theta, \mathcal{B}) + \sum_{\theta=1}^K \rho_i(\theta) \hat{U}_i(\theta; \theta, \mathcal{B}) \right),$$

which is equivalent to $\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}]$. Thus

$$Pr(\mathcal{E}) = \frac{C}{\mathbb{E}[\hat{\Psi}|\mathcal{B}] + \mathbb{E}[\hat{U}|\mathcal{B}] - 1}.$$

Substituting into (10) and dividing by C implies

$$\min \frac{\xi - (1 - \xi)\mathbb{E}[\widehat{U}|\mathcal{B}]}{\mathbb{E}[\widehat{\Psi}|\mathcal{B}] + \mathbb{E}[\widehat{U}|\mathcal{B}] - 1}.$$

The remaining constraints follow from plugging in for z_i using Theorem 1. Alternatively, we derive them from the Lagrangian in Appendix B. \square

A.13 Proof of Proposition 6

Proof. Applying Corollary 4, confidential arbitration means that the arbitrator can choose any distribution over consistent information structures. Each consistent information structure induces the belief system $B(s)$ which itself is consistent. Following much the same steps from the proof of Proposition 5 the arbitrator's objective becomes $\sum_s Pr(s)\mathcal{O}(s)$, where

$$\mathcal{O}(B(s)) := \frac{\xi + (1 - \xi) \left(1 - \mathbb{E}[\tilde{U}|B(s), s]\right)}{\mathbb{E}[\tilde{\Psi}|B(s), s] + \mathbb{E}[\tilde{U}|B(s), s] - 1}, \text{ with}$$

$$\rho_i(\theta_i|s) := Pr(\theta_i|\mathcal{E}, s),$$

$$\mathbb{E}[\tilde{\Psi}|B(s), s] := \sum_i \sum_{\theta \in \Theta} \rho_i(\iota_i(\theta_i)|s) \widehat{\Psi}_i(\theta_i, B(s)), \quad \text{and}$$

$$\mathbb{E}[\tilde{U}|B(s), s] := \sum_i \sum_{\theta \in \Theta} \rho_i(\theta|s) U(\theta; \theta, B(s)).$$

She minimizes this objective subject to the constraints stated in Problem (P_B) and subject to downward incentive constraints $(IC)^-$. If these additional constraints do not bind, it is easy to see that the optimal signal structure puts full mass on the belief system with the lowest value of \mathcal{O} , which proves the proposition.

If an additional constraint binds, the optimal signal can be found by setting up the Lagrangian function of the problem (see Lemma 10 in Appendix B) and choosing the signal distribution that concavifies the inverse of that function. \square

A.14 Proof of Proposition 7

Proof. Recall that $\gamma_i(m_i) = Pr(\mathcal{E})\rho_i(m_i)/p(m_i) = \sum_{\theta_{-i}} p(\theta_{-i})\gamma_i(\theta, \theta_{-i})$ is the expected probability of escalation given report m_i . We prove Proposition 7 as a special case of Lemma 8.

Lemma 8. *If $\gamma_i(m)D_i^+(m; k, \mathcal{B})$ is non-decreasing in m on some interval $[\underline{m}, \overline{m}]$ and $k \in [\underline{m}, \overline{m}]$, then local incentive compatibility for type k implies incentive com-*

patibility for any report in that interval.

Proof. Take k and m . Incentive compatibility holds iff

$$\begin{aligned} z_i(k) + \gamma_i(k)U_i(k; k, \mathcal{B}) &\geq z_i(m) + \gamma_i(m)U_i(m; k, \mathcal{B}) \\ \Leftrightarrow -\gamma_i(m)U_i(m; k, \mathcal{B}) &\geq z_i(m) - z_i(k) - \gamma_i(k)U_i(k; k, \mathcal{B}). \end{aligned} \quad (12)$$

Assume first that $m > k$. Adding and subtracting $\sum_{\theta=k+1}^m \gamma_i(m)U_i(m; \theta, \mathcal{B})$ to the LHS of (12) turns it into

$$\gamma_i(m) \left(\sum_{\theta=k}^{m-1} D_i^+(m; \theta, \mathcal{B}) - U_i(m; m, \mathcal{B}) \right).$$

Adding and subtracting $\sum_{\theta=k+1}^{m-1} z_i(\theta)$ to the RHS of (12) and using local downward incentive compatibility, i.e., $z_i(\theta) - z_i(\theta-1) \leq y_i(\theta-1, \theta-1) - y_i(\theta-1, \theta)$, implies

$$\sum_{\theta=k+1}^m (z_i(\theta) - z_i(\theta-1)) - \gamma_i(k)U_i(k; k, \mathcal{B}) \leq \sum_{\theta=k}^{m-1} \gamma_i(\theta-1)D_i^+(\theta-1; \theta, \mathcal{B}) - \gamma_i(m)U_i(m; m, \mathcal{B}).$$

The RHS of the above equation is an upper bound on the RHS of (12). Thus, (12) holds if

$$\sum_{\theta=k}^{m-1} \gamma_i(m)D_i^+(m; \theta, \mathcal{B}) \geq \sum_{\theta=k}^{m-1} \gamma_i(\theta-1)D_i^+(\theta-1; \theta, \mathcal{B}), \quad (13)$$

which holds since $\gamma_i(m)D_i^+(m; \theta, \mathcal{B})$ is non-decreasing in m .

For $m < k$, take equation (12), add and subtract $\sum_{\theta=m+1}^{k-1} \gamma_i(m)U_i(m; \theta, \mathcal{B})$ from the LHS. Iteratively applying local downward incentive compatibility to $z_i(m)$ and simplify to

$$\sum_{\theta=m}^{k-1} \gamma_i(\theta)D_i^+(\theta; \theta, \mathcal{B}) \geq \sum_{\theta=m+1}^{k-1} \gamma_i(m)D_i^+(m; \theta, \mathcal{B}),$$

which again holds since $\gamma_i(m)D_i^+(m; \theta, \mathcal{B})$ is non-decreasing in m . \square

The special case of Lemma 8 with $\underline{m} = 1$ and $\bar{m} = K$ for any k concludes the proof of Proposition 7. \square

A.15 Proof of Proposition 8

We prove Proposition 8 as a special case of Lemma 9

Lemma 9. *Local incentive constraints and (MDR) imply that $\gamma_i(m)D_i^+(m, k; \mathcal{B})$ is non-decreasing in m for any $m > k$.*

Proof. Take i and any m and $m-1$. Local incentive compatibility implies

$$y_i(m, m) - y_i(m, m-1) \geq z_i(m-1) - z(m) \geq y_i(m-1, m) - y_i(m-1, m-1).$$

Thus $\gamma_i(m)D_i^+(m; m, \mathcal{B}) \geq \gamma_i(m-1)D_i^+(m-1; m, \mathcal{B})$ or equivalently

$$\frac{\gamma_i(m)}{\gamma_i(m-1)} \geq \frac{D_i^+(m-1; m, \mathcal{B})}{D_i^+(m; m, \mathcal{B})}. \quad (14)$$

The term $\gamma_i(m)D_i^+(m; k, \mathcal{B})$ increases in m if $\gamma_i(m)D_i^+(m; k, \mathcal{B}) \geq \gamma_i(m-1)D_i^+(m-1; k, \mathcal{B})$ or equivalently,

$$\frac{\gamma_i(m)}{\gamma_i(m-1)} \geq \frac{D_i^+(m-1, k; \mathcal{B})}{D_i^+(m, k; \mathcal{B})} \quad (15)$$

which holds by (MDR) and (14) if $m > k$. \square

A.16 Proof of Proposition 9

Proof. Assume without loss that $\zeta(\theta_i) = 1/\theta_i$. A player's best-response to her opponent's action, $a(m_i; \theta_i, \mathcal{B})$, satisfies first-order conditions. The envelope theorem implies

$$\begin{aligned} U(m_i; \theta_i, \mathcal{B}) &= U(m_i; 1, \mathcal{B}) + \int_1^{\theta_i} c(a_i(m_i; s, \mathcal{B}))/s^2 ds \\ &= U(m_i; 1, \mathcal{B}) + h(m_i; \mathcal{B}) \int_1^{\theta_i} g(s) ds, \end{aligned}$$

where $g(s) := \tilde{g}(s)/s^2$ and where we used that $c(a_i(m_i; \theta_i, \mathcal{B})) = h(m_i; \mathcal{B})g(\theta_i)$. Thus, $D_i^+(m_i; \theta_i, \mathcal{B}) = h_i(m_i; \mathcal{B}) \int_{\theta_i}^{\theta_i+1} g(s) ds$. Moreover, for any m_i and m'_i we have that

$$\frac{D_i^+(m_i; \theta_i, \mathcal{B})}{D_i^+(m'_i; \theta_i, \mathcal{B})} = \frac{h_i(m_i; \mathcal{B}) \int_{\theta_i}^{\theta_i+1} g(s) ds}{h_i(m'_i; \mathcal{B}) \int_{\theta_i}^{\theta_i+1} g(s) ds} = \frac{h_i(m_i; \mathcal{B})}{h_i(m'_i; \mathcal{B})},$$

which is independent of θ_i . \square

B Lagrangian Problem

Remark. Our argument throughout this section assumes that $\bar{g}_B(K, K) = 1$. This normalization is without loss. For cases in which $0 < \bar{g}_B(K, K) < 1$ relabeling provides the missing step. The remaining cases with $\gamma(K, K) = 0$ are covered by continuity of B in γ . Lemma 7 in the proof of Theorem 1 provides the corresponding formal argument.

The designer's choice is $cs = (z, \gamma)$. The choice set is CS .

Lemma 10. *The Lagrangian approach yields the global optimum.*

Proof. We use Theorem 1 in Luenberger (1969) to show that the Lagrangian approach is sufficient. Let \mathbf{t} be the vector of Lagrangian multiplier. Further, let $G(\cdot)$ be the set of inequality constraints. Define $w(\mathbf{t}) := \inf\{-Obj|cs = (\gamma, z) \in CS, G(cs) \leq \mathbf{t}\}$, where Obj is the objective the designer wants to maximize. The Lagrangian is sufficient for a global optimum if $w(\mathbf{t})$ is convex.

Assume for a contradiction that $w(\mathbf{t}_0)$ is not convex at \mathbf{t}_0 . Then, there is $\mathbf{t}_1, \mathbf{t}_2$ and $x \in (0, 1)$ such that $x\mathbf{t}_1 + (1-x)\mathbf{t}_2 = \mathbf{t}_0$ and $xw(\mathbf{t}_1) + (1-x)w(\mathbf{t}_2) < w(\mathbf{t}_0)$. For $j \in \{1, 2\}$ let $cs_j = (\gamma[j], z[j])$ describe the optimal solution, such that $-Obj(cs_j) = w(\mathbf{t}_j)$. Note that $\gamma[j]$ induces $\mathcal{B}[j] = (B[j], S[j])$. Then, consider the choice cs_0 such that $z[0] = \lambda z[1] + (1-x)z[2]$, $\gamma[0] = x\gamma[1] + (1-x)\gamma[2]$ and $\Sigma[0] = \{\Sigma[1], \Sigma[2]\}$, with $Pr(\Sigma[1]) = x$. The choice cs_0 corresponds to a belief system $B[0]$ induced by some $\gamma[0]$. Moreover, it is common knowledge that $\Sigma[0]$ induces \mathcal{B} where realization $\mathcal{B}[1]$ occurs with probability x and $\mathcal{B}[2]$ with probability $1-x$. By construction all constraints are satisfied and the solution value equals that of the convex combination

$$w(\mathbf{t}_0) = -Obj(cs_0) = \sum_{j \in \{1, 2\}} Pr(\Sigma[j]) - Obj(\Sigma[j]) = xw(\mathbf{t}_1) + (1-x)w(\mathbf{t}_2).$$

A contradiction. □

For any i, θ , the constraints to the minimization problem are

$$\forall \theta \neq \theta' \quad -(z_i(\theta) - z_i(\theta')) - y_i(\theta; \theta) + y_i(\theta'; \theta) \leq 0, \quad (IC)$$

$$-z_i(\theta) - y_i(\theta; \theta) + V_i(\theta) \leq 0, \quad (PC_i)$$

$$-1 + \sum_i \sum_{\theta=1}^K p(\theta) z_i(\theta) + Pr(\mathcal{E}) \leq 0, \quad (RC)$$

$$\gamma(\theta_A, \theta_B) - 1 \leq 0. \quad (F)$$

We now derive the Lagrangian representation of the optimization problem. First, we state the complementary slackness conditions and the respective Lagrangian

multipliers

$$\begin{aligned}
[z_i(\theta) - z_i(\theta') + y_i(\theta; \theta) - y_i(\theta'; \theta)]\nu_{\theta, \theta'}^i &= 0, & \nu_{\theta, \theta'}^i &\geq 0; \\
[z_i(\theta) + y_i(\theta; \theta) - V_i(\theta)]\lambda_\theta^i &= 0, & \lambda_\theta^i &\geq 0; \\
\left[1 - \sum_i \sum_\theta p(\theta) z_i(\theta) - \Pr(\mathcal{E})\right] \delta &= 0, & \delta &\geq 0; \\
[1 - \gamma(\theta_A, \theta_B)]\mu_{\theta_A, \theta_B} &= 0, & \mu_{\theta_A, \theta_B} &\geq 0.
\end{aligned}$$

For any Lagrangian multiplier, say t , we introduce the following notation $\tilde{t} \equiv \frac{t}{\delta}$. Define

$$\tilde{\Lambda}^i(\theta) := \sum_{k=1}^{\theta} \tilde{\lambda}_k^i. \quad (16)$$

Next, we characterize the solution in terms of the Lagrangian objective.

Lemma 11. *\mathcal{B} is an optimal solution to the designers problem if and only if there are Lagrangian multipliers that satisfy complementary slackness and \mathcal{B} maximizes*

$$\frac{(1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}] - 1}{\hat{\mathcal{L}}(\mathcal{B}) - 1},$$

$$\begin{aligned}
\text{where } \hat{\mathcal{L}}(\mathcal{B}) := & \mathcal{T}(\mathcal{B}) + \sum_i \left[\sum_{\theta=1}^K \rho_i(\theta) \hat{U}_i(\theta; \theta, \mathcal{B}) \right. \\
& + \sum_{\theta=1}^{K-1} \sum_{\theta'=\theta+1}^K \frac{\tilde{\nu}_{\theta, \theta'}^i + \mathbf{M}^i(\theta) - \tilde{\nu}^i(\theta, \theta')}{p(\theta)} \rho_i(\theta) \{ \hat{U}_i(\theta; \theta, \mathcal{B}) - \hat{U}_i(\theta; \theta', \mathcal{B}) \} \\
& \left. - \sum_{\theta=1}^{K-1} \sum_{\theta'=\theta+1}^K \frac{\tilde{\nu}_{\theta, \theta'}^i}{p(\theta')} \rho_i(\theta') \{ \hat{U}_i(\theta'; \theta, \mathcal{B}) - \hat{U}_i(\theta'; \theta', \mathcal{B}) \} \right], \quad (17)
\end{aligned}$$

where $\mathbf{M}^i(\theta) := \tilde{\Lambda}^i(\theta) - \sum_{k=1}^{\theta} p(k)$, and $\tilde{\nu}_i(\theta, \theta') := (\sum_{k=1}^{\theta} \sum_{\tilde{\theta} > \theta}^K (\tilde{\nu}_{\tilde{\theta}, k}^i - \tilde{\nu}_{k, \tilde{\theta}}^i)) - (\tilde{\nu}_{\theta', \theta}^i - \tilde{\nu}_{\theta, \theta'}^i)$.

$$\mathcal{T}(\mathcal{B}) := - \sum_{\theta_A \times \theta_B} \frac{\rho_A(\theta_A) \beta_A(\theta_B | \theta_A)}{p(\theta_A) p(\theta_B)} \tilde{\mu}_{\theta_A, \theta_B}. \quad (18)$$

Moreover, the following is true at the optimum:

- Constraint (BB) is always binding, i.e., $\delta > 0$.
- $\mathbf{M}^i(\theta) = \tilde{\nu}^i(\theta, \theta') + \tilde{\nu}_{\theta', \theta}^i - \tilde{\nu}_{\theta, \theta'}^i$ for any θ' .
- If $\tilde{\Lambda}^i(\theta) - \sum_{j=1}^{\theta} p(j) > 0$, then there is at least one type $k \leq \theta$ such that this

type's downward incentive constraint is binding. If in addition the upward incentive constraints are redundant, then $\tilde{\nu}_{\theta,\theta'}^i = 0$ for all $\theta' \geq \theta$.

- If $\tilde{\Lambda}^i(\theta) - \sum_{j=1}^{\theta} p(j) < 0$, then there is at least one type $k \leq \theta$ such that this type's upward incentive constraint is binding. If in addition the downward incentive constraints are redundant, then $\tilde{\nu}_{\theta,\theta'}^i = 0$ for all $\theta' < \theta$.
- If local incentive constraints are sufficient, then $\tilde{\nu}_{\theta,\theta'}^i = 0$ for any θ such that $\theta' > \theta+1$ or $\theta' < \theta-1$. Moreover, $\tilde{\nu}^i(\theta, \theta') = \mathbf{M}^i(\theta)$ for any θ, θ' such that $\theta' \neq \{\theta-1, \theta+1\}$.

Proof. We manipulate the Lagrangian, \mathcal{L} , and derive a more tractable dual problem. We want to minimize $\xi \Pr(\mathcal{E}) + (1-\xi) \left(-\sum_i \sum_{\theta=1}^K p(\theta) [z_i(\theta) + \gamma_i(\theta) \hat{U}_i(\theta; \theta, \mathcal{B})] \right)$. We first relax the problem by replacing $\sum_i \sum_{\theta=1}^K p(\theta) z_i(\theta)$ with $1 - \Pr(\mathcal{E})$. Factoring out $\Pr(\mathcal{E})$ from the objective and applying Bayes rule, the objective becomes $\Pr(\mathcal{E})(\xi - (1-\xi)\mathbb{E}[\hat{U}|\mathcal{B}]) - (1-\xi)$. Dropping the constant $(1-\xi)$, the Lagrangian reads

$$\begin{aligned} \mathcal{L} = & \Pr(\mathcal{E}) \left(1 - (1-\xi)\mathbb{E}[\hat{U}|\mathcal{B}] \right) + \delta [-1 + \sum_i \sum_{\theta=1}^K p(\theta) z_i(\theta) + \Pr(\mathcal{E})] \\ & + \sum_i \sum_{\theta=1}^K [-z_i(\theta) - y_i(\theta; \theta) + V_i(\theta)] \lambda_{\theta}^i \\ & + \sum_i \sum_{\theta=1}^K \sum_{\theta' \in \Theta \setminus \theta} [-z_i(\theta) + z_i(\theta') - y_i(\theta; \theta) + y_i(\theta'; \theta)] \nu_{\theta,\theta'}^i \\ & + \sum_{\theta_A \times \theta_B} [\gamma(\theta_A, \theta_B) - 1] \mu_{\theta_A, \theta_B}. \end{aligned} \tag{19}$$

Using Theorem 1 and Lemma 3 we optimize over $\{z_i(\cdot), \gamma(K, K), \{\beta_A(\cdot|j)\}_{j \in \Theta \cup \beta_B(\cdot|1)}\}$, with $\gamma(K, K) := \Pr(\mathcal{E}|\theta_A=K, \theta_B=K)$.

Step 1: Eliminating $z_i(\cdot)$ using First-order Conditions. Define $\nu_{K+1,K}^i := 0 =: \nu_{1,0}^i = \nu_{0,1}^i$ for ease of notation. The FOC w.r.t. $z_i(\theta)$ are

$$p(\theta)\delta - \lambda_{\theta}^i + \sum_{\tilde{\theta} \in \Theta_i \setminus \theta} (\nu_{\tilde{\theta},\theta}^i - \nu_{\theta,\tilde{\theta}}^i) = 0. \tag{20}$$

Summing over all K conditions in (20) and recalling definition (16) yields

$$1 = \tilde{\Lambda}^i(K). \tag{21}$$

(20) holds for all θ if and only if

$$\sum_{k=1}^{\theta} \sum_{\tilde{\theta} > \theta} (\tilde{\nu}_{\tilde{\theta},k}^i - \tilde{\nu}_{k,\tilde{\theta}}^i) = - \sum_{k=1}^{\theta} p(k) + \tilde{\Lambda}^i(\theta). \tag{22}$$

Thus, $\sum_{k=1}^{\theta} \sum_{\hat{\theta} > \theta}^K \tilde{\nu}_{\hat{\theta},k}^i > 0$ if $M^i(\theta) > 0$ and vice versa for $\sum_{k=1}^{\theta} \sum_{\hat{\theta} > \theta}^K \tilde{\nu}_{k,\hat{\theta}}^i$. We solve (20) for $\tilde{\lambda}_{\hat{\theta}}^i$ and substitute into (19). We also substitute $\tilde{\nu}_{\hat{\theta}',\theta}^i = M^i(\theta) + \tilde{\nu}^i(\theta, \theta') - \tilde{\nu}_{\hat{\theta}',\theta}^i$ for all $\theta' > \theta$ into (19) and sort terms. Moreover, all terms involving $z_i(\cdot)$ cancel out from (19) via (20).

Step 2: Reformulating the Lagrangian Objective. Given the above necessary conditions, we manipulate the Lagrangian objective to derive a more tractable maximization problem. Using Bayes' rule together with the homogeneity established in the proof of Theorem 1 (step 1), applying algebra and using the first-order conditions it is straightforward to show that (19) admits the following representation

$$\mathcal{L} = Pr(\mathcal{E})(\delta + \xi - (1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}]) - \delta C - \delta \sum_{\sigma} Pr(\mathcal{E})\hat{\mathcal{L}}(\mathcal{B}), \quad (23)$$

where C is a constant that is independent of the choice variables and reads

$$C := 1 - \sum_i \sum_{\theta} \tilde{\lambda}_{\theta}^i V_i(\theta) + \sum_{\theta_A \times \theta_B} \tilde{\mu}_{\theta_A, \theta_B} < 0.$$

Define $\gamma(\theta_A, \theta_B) := Pr(\mathcal{E}|\theta_A, \theta_B)$. From the proof of Theorem 1 (step 1) with $\alpha = \gamma(K, K)$ it follows that $\gamma(\theta_A, \theta_B) = f(\mathcal{B}, \theta_A, \theta_B)\gamma(K, K)$, where $f(\mathcal{B}, \theta_A, \theta_B)$ is a positive real number. Thus, $Pr(\mathcal{E}) = \gamma(K, K)R(\mathcal{B})$ with $R(\mathcal{B}) := \sum_{\theta_A \times \theta_B} p(\theta_A)p(\theta_B)f(\mathcal{B}, \theta_A, \theta_B)$. Plugging into (23) yields

$$\mathcal{L} = \gamma(K, K)R(\mathcal{B})(\delta - (1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}]) - \delta C - \delta \gamma(K, K)R(\mathcal{B})\hat{\mathcal{L}}(\mathcal{B}). \quad (24)$$

The FOC of (24) w.r.t. $\gamma(K, K)$ is

$$R(\mathcal{B})\left(\delta + 1 - (1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}] - \delta\hat{\mathcal{L}}(\mathcal{B})\right) = 0. \quad (25)$$

By Assumption 3 $R(\mathcal{B}) > 0$ and thus, $\hat{\mathcal{L}}(\mathcal{B}) - 1 > 0$ if $\gamma(K, K) > 0$. Therefore, $\delta = (1 - (1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}])(\hat{\mathcal{L}}(\mathcal{B}) - 1)^{-1}$. Substituting into (24) and simplifying yields

$$\mathcal{L} = \frac{(-C)(1 - (1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}])}{\hat{\mathcal{L}}(\mathcal{B}) - 1}, \quad (26)$$

which is minimized if and only if $\frac{(1 - \xi)\mathbb{E}[\hat{U}|\mathcal{B}] - 1}{\hat{\mathcal{L}}(\mathcal{B}) - 1}$ is maximized.²⁰

□

²⁰The Lagrangian multipliers are such that C is negative at the optimum. Otherwise (19) and (26) imply that $Pr(\mathcal{E})$ is negative, a contradiction to Assumption 3 or to existence of the following mechanism. Take a degenerate signal distribution and set $\gamma(\theta_A, \theta_B) = 1$ for all type profiles.

C Example: Type-Dependent Lotteries

In this section we provide an algorithm for the conflict minimization in Example 2. We restrict the environment to be monotone and focus on constant surplus reduction in case of conflict. That makes ξ irrelevant for our problem. This setting nests the solutions from the literature such as Fey and Ramsay (2011), Hörner, Morelli, and Squintani (2015), and the monotone cases in Bester and Wärneryd (2006). We comment on changes when relaxing these assumptions at the end.

Definition 9 (Lottery). \mathcal{E} is a lottery if $U_i(\theta_i; \theta_i, \mathcal{B} | \theta_{-i})$ is constant in \mathcal{B} . The lottery is *0-sum* if $u(\theta_i, \theta_{-i}) + u(\theta_{-i}, \theta_i)$ is constant for all (θ_i, θ_{-i}) .

The problem is linear in the joint distribution over type-pairs in the event \mathcal{E} , ρ . There is a one-to-one relationship to (consistent) B . Due to the linearity we can ignore signals. We abuse notation and replace \mathcal{B} by the function ρ . To simplify, we assume $V_i(\theta_i) = \sum_i p(\theta_{-i})u(\theta_i, \theta_{-i})$ and impose monotonicity.

Definition 10 (Monotone Lottery). A lottery is monotone if

- $u(\theta_i, \theta_{-i}) - u(\theta_i - 1, \theta_{-i})$ is weakly increasing in θ_i and θ_{-i} , and
- the prior p induces a weakly increasing inverse hazard rate $\omega(\theta_i) := \frac{\sum_{k=1}^{\theta_i} p(k)}{p(\theta_i)}$.

In monotone lotteries $\iota_i(\theta_i) = \theta_i + 1$, i.e., local (upward) incentive constraints imply incentive compatibility. Define

$$\Upsilon(\theta_A, \theta_B) := \omega(\theta_A) (u(\theta_A, \theta_B) - u(\theta_A - 1, \theta_B)) + \omega(\theta_B) (u(\theta_B, \theta_A) - u(\theta_B - 1, \theta_A)).$$

Restricting attention (without loss) to the $\xi = 1$ case and plugging into the objective yields

$$2u(K, K) + \sum_{\Theta^2 \setminus (1,1)} \rho(\theta_A, \theta_B) \Upsilon(\theta_A, \theta_B) =: \mathcal{O}(\rho),$$

the objective the arbitrator wishes to maximize. Identifying the highest Υ and setting the corresponding ρ equal to 1 achieves that. If $M(\rho) \neq 0$ for that ρ the problem is solved. Otherwise the optimal solution is not feasible. Instead, the arbitrator has to increase available funds by putting some weight on the second highest Υ as well. For monotone 0-sum lotteries there is a simple algorithm to construct optimal arbitration.

Definition 11 (Top-Down Algorithm). Let Θ_+^2 be the set of type pairs (θ_A, θ_B) such that $\rho(\theta_A, \theta_B) > 0$. Begin by setting $\Theta_+^2 = \emptyset$.

1. Set $\rho(K, K) = 1$ and check if $\rho(K, K) \leq \frac{(p(K))^2(\mathcal{O}(\rho)-1)}{2V(K)-1}$. If it holds, terminate. Otherwise continue at 2.

2. Identify the set $\Theta_N^2 = \{(\theta_A, \theta_B) | (\theta_A, \theta_B) = \arg \max_{\Theta^2 \setminus \Theta_+^2} \Upsilon(\theta_A, \theta_B)\}$.

(a) Set $\rho(K, K)$ to the solution of

$$\sum_{\substack{(\theta_A, \theta_B) \in \\ \Theta_+^2 \cup \Theta_N^2}} \frac{p(\theta_A)p(\theta_B)}{(p(K))^2} \rho(K, K) = 1. \quad (27)$$

(b) Replace $\rho(\theta_A, \theta_B) = \frac{p(\theta_A)p(\theta_B)}{(p(K))^2} \rho(K, K) \forall (\theta_A, \theta_B) \in \Theta_+^2 \cup \Theta_N^2$.

(c) Check whether the condition in 1 holds. If it holds, decrease all ρ for the set Θ_N^2 at the expense $\rho(K, K)$ keeping the relation of 2(b) until the condition holds with equality. Then, terminate. If it is violated, repeat step 2.

Proposition 10. *Suppose the escalation game is a monotone 0-sum lottery. Optimal arbitration is the solution to the top-down algorithm.*

Proof. Jointly conditions from step 1 and equation (27) are necessary and sufficient for (BB).²¹

By construction the top-down algorithm point-wise maximizes $\mathcal{O}(\rho)$ subject to (BB). What remains is to show that all ignored constraints are satisfied. We show this using monotonicity, i.e., $\gamma(\theta_i+1, \theta_{-i}) \geq \gamma(\theta_i, \theta_{-i}) \Leftrightarrow p(\theta_i)\rho(\theta_i+1, \theta_{-i}) \geq p(\theta_i+1)\rho(\theta_i, \theta_{-i})$ for all θ_i, θ_{-i} .

Monotonicity trivially holds if $\gamma(K, K) \neq 1$ because it implies $\rho(K, K) = 1$. Thus, assume $\gamma(K, K) = 1$. By Bayes' rule

$$\gamma(\theta_A, \theta_B) = \left(\rho(\theta_A, \theta_B) \right) / \left(p(\theta_A)p(\theta_B) \right) Pr(\mathcal{E}).$$

When $\gamma(\theta_A, \theta_B) > 0$, then $\gamma(\theta_A+1, \theta_B) = 1$ and

$$Pr(\mathcal{E})\rho(\theta_A+1, \theta_A) = p(\theta_A+1)p(\theta_B), \quad Pr(\mathcal{E})\rho(\theta_A, \theta_A) \leq p(\theta_A)p(\theta_B).$$

Monotonicity holds since $p(\theta_A)\rho(\theta_A+1, \theta_A) \geq p(\theta_A+1)\rho(\theta_A, \theta_A)$, and all but the local upward incentive constraints are redundant because

$$\sum_{\theta_{-i}=1}^K (p(\theta'_i)\rho(\theta_i, \theta_{-i}) - p(\theta_i)\rho(\theta'_i, \theta_{-i})) [u(\theta_i, \theta_{-i}) - u(\theta'_i, \theta_{-i})] \geq 0. \quad (28)$$

Finally, we verify that only the highest type's participation constraint binds at the optimum. It implies that upward local incentive constraints hold with equality. We

²¹Substitute $\gamma(\theta_A, \theta_B)p(\theta_A)p(\theta_B)(Pr(\mathcal{E}))^{-1}$ for $\rho(\theta_A, \theta_B)$ in the RHS of (BB) and substitute in the LHS accordingly using the path ι .

verify the claim by induction. We first show that $\Pi_i(K-1; K-1) \geq V(K-1)$. By local incentive compatibility and $\Pi_i(K; K) \geq V(K)$,²² we know that $\Pi_i(K-1; K-1) \geq \Pi_i(K, K) - y_i(K; K) + y_i(K; K-1)$. Thus, to show that $\Pi_i(K-1; K-1) - V(K-1) \geq 0$ it suffices to show that $\Pi_i(K, K) - V(K-1) \geq y_i(K; K) - y_i(K; K-1)$. The game is a lottery, and we need

$$\sum_{\theta_{-i}} (p(\theta_{-i}) - \gamma_i(K)\beta_i(\theta_{-i}|K)) (u(K, \theta_{-i}) - u(K-1, \theta_{-i})) \geq 0.$$

The last bracket is positive by monotonicity. The first is $(p(\theta_{-i}) - \gamma_i(K)\beta_i(\theta_{-i}|K)) = (Pr(K, \theta_{-i}) - Pr(K, \theta_{-i}, \mathcal{E})) / p(K) \geq 0$.

None of the above argument depends on the specifics of K , thus the induction step to verify $\Pi_i(\theta_i; \theta_i) \geq V(\theta_i)$ for all θ_i follows analogously. \square

If we give up the 0-sum element of the lottery, welfare maximization is no longer isomorphic to escalation minimization. Let $W(\theta_A, \theta_B) = u(\theta_A, \theta_B) + u(\theta_A, \theta_B)$ be joint surplus of a type pair (θ_A, θ_B) . Define

$$\mathcal{O}'(\rho) := 2u(K, K) + \sum_{\Theta^2 \setminus (1,1)} \rho(\theta_A, \theta_B) (\Upsilon(\theta_A, \theta_B) + W(\theta_A, \theta_B) - W(1, 1))$$

The general problem becomes

$$\min_{\rho} \frac{\xi + (1 - \xi) \sum_{\Theta^2} \rho(\theta_A, \theta_B) W(\theta_A, \theta_B)}{\mathcal{O}'(\rho) - 1},$$

$$\text{s.t. } \rho(K, K) \leq \frac{(p(K))^2 (\mathcal{O}'(\rho) - 1)}{2V(K) - 1} \text{ and } \sum_{\substack{(\theta_A, \theta_B) \in \\ \Theta_+^2 \cup \Theta_N^2}} \frac{p(\theta_A)p(\theta_B)}{(p(K))^2} \rho(K, K) = 1.$$

Without 0-sum and $\xi < 1$ the objective is not linear and thus harder to solve. Yet, it remains that signals are of no help because U remains linear in the information structure leaving no room to exploit the curvature.

D Solution Algorithm for Monotone Mechanisms

Here we provide details behind the results obtained for type-separable escalation games at the end of Section 5. Recall that these escalation games feature the following payoff structure.

$$u(a_i; a_{-i}, \theta_i) = \phi(a_i, a_{-i}) - \zeta(\theta_i)c(a_i, a_{-i}). \quad (29)$$

²²In the optimal mechanism it holds that $\Pi_i(K; K) = V(K)$.

Let a_{-i}^* be the equilibrium action of player $-i$ and define $c(a_i) := \mathbb{E}[c(a_i, a_{-i}^*) | m_i, \mathcal{B}]$.

First, we derive the designer's objective if ϕ only distributes the pie without destroying any surplus.

The envelope theorem implies

$$U_i(m_i; \theta_i, \mathcal{B}) = \mathbb{E}[\phi(a_i(m_i; 1, \mathcal{B}), a_{-i}^*) | m_i, \mathcal{B}] - C(m_i, \theta_i; \mathcal{B}), \text{ where} \quad (30)$$

where $C(m_i; \theta_i, \mathcal{B}) := \int_{\zeta(1)}^{\zeta(\theta_i)} c(m_i; s, \mathcal{B}) d(-\zeta(s)) + \zeta(1)c(m_i; 1, \mathcal{B})$ and let $c(m_i; s, \mathcal{B}) := c(a_i^*(m_i; \theta_i, \mathcal{B}))$ be θ_i 's expected cost from his optimal action $a_i^*(m_i; \theta_i, \mathcal{B})$. If $\xi = 1$, optimal arbitration maximizes $\sum_i (\mathbb{E}[U_i | \mathcal{B}] + \mathbb{E}[\Psi_i | \mathcal{B}])$. Since $\phi(a_i, a_{-i}) + \phi(a_{-i}, a_i) = 1$, $\sum_i \mathbb{E}[U_i | \mathcal{B}] = 1 - \sum_i \rho_i(\theta_i) \zeta(\theta_i) c(\theta_i; \theta_i, \mathcal{B})$. Moreover, (30) implies that $D_i^+(m_i; \theta_i - 1, \mathcal{B}) = -C(m_i; \theta_i, \mathcal{B}) + C(m_i; \theta_i - 1, \mathcal{B})$.

As we will show below, in the optimum upward adjacent incentive constraints imply incentive compatibility. Thus, $\iota_i(\theta) = \theta + 1$, $w_i^t(\theta) = \sum_{k=1}^{\theta+1} p(k)/p(\theta+1) =: \omega(\theta+1)$, and $D_i^c = D_i^+$. Let

$$\tilde{S}(\theta_i, \theta_{-i}; \mathcal{B}) := \sum_i \omega(\theta_i) D_i^+(\theta_i; \theta_i - 1, \mathcal{B}) - \zeta(\theta_i) c(\theta_i; \theta_i, \mathcal{B}).$$

The objective becomes $\mathcal{O}(\rho(\cdot, \cdot)) := \sum \rho(\theta_1, \theta_2) \tilde{S}(\theta_1, \theta_2; \mathcal{B})$. If the type space is sufficiently dense, we can set up an auxiliary problem. We replace \tilde{S} with S being defined as

$$S(\theta_i, \theta_{-i}; \mathcal{B}) := (\omega(\theta_i) \Delta \theta - \zeta(\theta_i)) c(\theta_i; \theta_i, \mathcal{B}) + (\omega(\theta_{-i}) - \zeta(\theta_{-i})) c_{-i}(\theta_{-i}; \theta_{-i}, \mathcal{B}).$$

Sufficiency of the Auxiliary Problem. We show that the solution to the auxiliary problem solves the original problem if $\Delta \theta$, for any two adjacent types, is sufficiently small. By the intermediate value theorem we have that

$$\tilde{S}(\theta_i, \theta_{-i}; \mathcal{B}) = \sum_i \omega(\theta_i) \Delta \theta c(\theta_i; \tilde{\theta}_i, \mathcal{B}) - \zeta(\theta_i) c(\theta_i; \theta_i, \mathcal{B})$$

for some $\tilde{\theta}_i \in [\theta_i - 1, \theta_i]$. The objective becomes $\sum_i \tilde{\omega}(\theta_i) c(\theta_i; \tilde{\theta}_i, \mathcal{B}) - \theta_i c(\theta_i; \theta_i, \mathcal{B})$, where $\tilde{\omega} := \omega \Delta \theta$. If $\Delta \rightarrow 0$, then $c(\theta_i; \tilde{\theta}_i, \mathcal{B}) \rightarrow c(\theta_i; \theta_i, \mathcal{B})$ and the scores of the auxiliary problem and the original problem coincide. Hence the solutions coincide.

Results. Suppose ω is non-decreasing in θ . Moreover, assume the distance between any two adjacent types, $\Delta \theta := \zeta(\theta - 1) - \zeta(\theta)$, is sufficiently small. We will show that, if the game features strategic complements, upward adjacent incentive constraints are necessary and sufficient for all other constraints. In particular, (4) is satisfied. Now, we state an algorithm that solves the problem.

We use the general Lagrangian approach from appendix B to develop a solution algorithm for our class of games. We apply it to the the auxiliary problem. We first relax that problem by ignoring all global incentive constraints. Then, the Lagrangian of the reduced-form problem becomes

$$\sum_{\theta_i \times \theta_{-i}} \rho(\theta_i, \theta_{-i}) \left(\left(\sum_{j \in \{i, -i\}} (\tilde{\omega}(\theta_j) - \zeta(\theta_j)) c_j(\theta_j, \theta_j; \mathcal{B}) \right) - \frac{\mu(\theta_i, \theta_{-i})}{p(\theta_i) p(\theta_{-i})} \right), \quad (31)$$

where $\mu(\theta_i, \theta_{-i})$ is the Lagrangian multiplier on the feasibility constraints, i.e.,

$$2V(1)\rho(K, K) \leq (p(K))^2 (\mathcal{O}(\rho(\cdot, \cdot))). \quad (32)$$

If that constraint does not bind, then the optimal solution features $\rho(K, K) = 1$. This follows from the complementary nature of the conflict, together with the non-decreasing virtual valuations.

Assume that $\rho(K, K) = 1$ is not feasible, that is, $\mu(K, K) > 0$. Then, the least-constrained solution is not feasible and signals may improve.

We state an algorithm, a top-down version with information revelation, and then argue that this algorithm is optimal.

Algorithm. Define the score of a type profile the following way:

$$\hat{S}(\theta_i, \theta_{-i}) = (\omega(\theta_i) - \zeta(\theta_i)) c_i(\theta_i; \theta_i, \mathcal{B}^{\theta_i, \theta_{-i}}) + (\omega(\theta_{-i}) - \zeta(\theta_{-i})) c_{-i}(\theta_{-i}; \theta_{-i}, \mathcal{B}^{\theta_i, \theta_{-i}}),$$

where $\mathcal{B}^{\theta_i, \theta_{-i}}$ is the belief system that results if each match receives full information. We order type profiles according to their score. If the highest type profile is (θ_i, θ_{-i}) , then the next highest type profile is either $(\theta_i - 1, \theta_{-i})$ or $(\theta_i, \theta_{-i} - 1)$ by complements.

Signals might improve because the least-constrained problem is not feasible. Hence given active type profiles, implement the optimal information revelation policy, i.e., that which maximizes the objective given the active type profiles. Check whether the objective satisfies constraint (32). If not, continue to the next highest type profile and repeat the maximization.

Optimality of the Algorithm. The increasing hazard rate and strategic complements imply that higher type profiles have a higher score. Moreover, strategic complements imply that, given the optimal information disclosure policy, type profiles with the highest ex-post scores are the most beneficial ones.

Optimal Information Revelation. To construct the concave hull of the Lagrangian objective, (31), we have to distinguish two cases. Define $c(a_i^*(a_{-i}))$ as that part of the cost that is independent of $-i$'s action. If that function is concave, then

we disclose no information. In contrast, if that function is convex, then we disclose full information.

Secondly, observe that the form of the Lagrangian objective (31) implies that given $\rho(\cdot, \cdot)$ the information disclosure that maximizes the least-constrained objective is optimal.

Incentive Constraints. We need to verify that this algorithm satisfies the incentive constraints. That is, $\gamma_i(m_i)D_i^+(m_i; \theta_i, \mathcal{B})$ is weakly increasing in m_i . Both $D_i^+(m_i; \theta_i, \mathcal{B}) = c_i(\theta_i, \theta_i; \mathcal{B})$ and $\gamma_i(\theta_i)$ are weakly increasing in m_i by complementarities and the structure of the optimal solution.

References

- Armstrong, M. J. and W. Hurley (2002). “Arbitration using the closest offer principle of arbitrator behavior”. *Mathematical Social Sciences* 43, pp. 19–26.
- Atakan, A. E. and M. Ekmekci (2014). “Auctions, actions, and the failure of information aggregation”. *American Economic Review* 104, pp. 2014–2048.
- Aumann, R. J. and M. Maschler (1995). *Repeated games with incomplete information*. Cambridge, MA: MIT press.
- Balzer, B. and J. Schneider (2017). “Optimal Alternative Dispute Resolution”. *mimeo*.
- (2018). “Persuading to Participate: Coordinating on a Standard”. *mimeo*.
- Bebchuk, L. A. (1984). “Litigation and settlement under imperfect information”. *The RAND Journal of Economics*, pp. 404–415.
- Bergemann, D. and S. Morris (2016). “Bayes correlated equilibrium and the comparison of information structures”. *Theoretical Economics*, pp. 487–522.
- Bester, H. and K. Wärneryd (2006). “Conflict and the Social Contract”. *The Scandinavian Journal of Economics* 108, pp. 231–249.
- Border, K. C. (2007). “Reduced form auctions revisited”. *Economic Theory* 31, pp. 167–181.
- Börger, T. and P. Norman (2009). “A note on budget balance under interim participation constraints: the case of independent types”. *Economic Theory* 39, pp. 477–489.
- Brown, J. G. and I. Ayres (1994). “Economic rationales for mediation”. *Virginia Law Review*, pp. 323–402.
- Calzolari, G. and A. Pavan (2006a). “Monopoly with Resale”. *The RAND Journal of Economics* 37, pp. 362–375.

- Calzolari, G. and A. Pavan (2006b). “On the optimality of privacy in sequential contracting”. *Journal of Economic theory* 130, pp. 168–204.
- Carroll, G. D. and I. R. Segal (2018). “Robustly Optimal Auctions with Unknown Resale Opportunities”. *mimeo*.
- Celik, G. and M. Peters (2011). “Equilibrium rejection of a mechanism”. *Games and Economic Behavior* 73, pp. 375–387.
- Compte, O. and P. Jehiel (2009). “Veto constraint in mechanism design: inefficiency with correlated types”. *American Economic Journal: Microeconomics* 1, pp. 182–206.
- Cr  mer, J. and R. P. McLean (1985). “Optimal Selling Strategies under Uncertainty for a Discriminating Monopolist when Demands are Interdependent”. *Econometrica* 53, pp. 345–361.
- (1988). “Full Extraction of the Surplus in Bayesian and Dominant Strategy Auctions”. *Econometrica* 56, pp. 1247–1257.
- Doornik, K. (2014). “A rationale for mediation and its optimal use”. *International Review of Law and Economics* 38, pp. 1–10.
- Doval, L. and V. Skreta (2018). “Constrained information design: Toolkit”. *arXiv preprint arXiv:1811.03588*.
- Dworczak, P. (2017). “Mechanism Design with Aftermarkets: Cutoff Mechanisms”. *mimeo*.
- Dworczak, P. and G. Martini (2018). “The simple economics of optimal persuasion”. *mimeo*.
- Fey, M. and K. W. Ramsay (2011). “Uncertainty and Incentives in Crisis Bargaining: Game-Free Analysis of International Conflict”. *American Journal of Political Science* 55, pp. 149–169.
- Fudenberg, D. and J. Tirole (1988). “Perfect Bayesian and Sequential Equilibria - A clarifying Note”. *mimeo*.
- Galperti, S. and J. Prego (2018). *A Dual Perspective on Information Design*. Tech. rep.
- Georgiadis, G. and B. Szentes (2018). “Optimal monitoring design”. *mimeo*.
- Goeree, J. K. (2003). “Bidding for the future: Signaling in auctions with an after-market”. *Journal of Economic Theory* 108, pp. 345–364.
- Goltsman, M., J. H  rner, G. Pavlov, and F. Squintani (2009). “Mediation, arbitration and negotiation”. *Journal of Economic Theory* 144, pp. 1397–1420.
- Gupta, M. and B. Lebrun (1999). “First price auctions with resale”. *Economics Letters* 64, pp. 181–185.
- H  rner, J., M. Morelli, and F. Squintani (2015). “Mediation and Peace”. *The Review of Economic Studies* 82, pp. 1483–1501.
- Jackson, M. O. and M. Morelli (2011). “The reasons for wars: an updated survey”. *The handbook on the political economy of war* 34.

- Jehiel, P. and B. Moldovanu (2001). “Efficient design with interdependent valuations”. *Econometrica* 69, pp. 1237–1259.
- Kamenica, E. and M. Gentzkow (2011). “Bayesian Persuasion”. *American Economic Review* 101, pp. 2590–2615.
- Kolotilin, A. (2018). “Optimal information disclosure: A linear programming approach”. *Theoretical Economics* 13, pp. 607–635.
- Kolotilin, A. and A. Zapechelnyuk (2019). “Persuasion meets delegation”. *mimeo*.
- Lauermann, S. and G. Virág (2012). “Auctions in Markets: Common Outside Options and the Continuation Value Effect”. *American Economic Journal: Microeconomics* 4, pp. 107–130.
- Luenberger, D. G. (1969). *Optimization by Vector Space Methods*.
- Mathevet, L., J. Perego, and I. Taneva (2017). “On Information Design in Games”. *mimeo*.
- Meirowitz, A., M. Morelli, K. W. Ramsay, and F. Squintani (2017). “Dispute Resolution Institutions and Strategic Militarization”. *Journal of Political Economy*, forthcoming.
- Myerson, R. B. (1982). “Optimal coordination mechanisms in generalized principal–agent problems”. *Journal of Mathematical Economics* 10, pp. 67–81.
- (1981). “Optimal Auction Design”. *Mathematics of Operations Research* 6, pp. 58–73.
- Myerson, R. B. and M. A. Satterthwaite (1983). “Efficient mechanisms for bilateral trading”. *Journal of Economic Theory* 29, pp. 265–281.
- Mylovanov, T. and A. Zapechelnyuk (2013). “Optimal arbitration”. *International Economic Review* 54, pp. 769–785.
- Olszewski, W. (2011). “A Welfare Analysis of Arbitration”. *American Economic Journal: Microeconomics* 3, pp. 174–213.
- Pavan, A. and G. Calzolari (2009). “Sequential contracting with multiple principals”. *Journal of Economic Theory* 144, pp. 503–531.
- Schweizer, U. (1989). “Litigation and settlement under two-sided incomplete information”. *The Review of Economic Studies* 56, pp. 163–177.
- Spier, K. E. (1994). “Pretrial bargaining and the design of fee-shifting rules”. *The RAND Journal of Economics*, pp. 197–214.
- Sugaya, T. and A. Wolitzky (2018). “Revelation Principles in Multistage Games”. *mimeo*.
- Zhang, J. (2014). “Optimal Mechanism Design with Aftermarket Interactions”. *mimeo*.
- Zheng, C. (2002). “Optimal Auction with Resale”. *Econometrica* 70, pp. 2197–2224.
- (2018). “A necessary and sufficient condition for Peace”. *mimeo*.

Zheng, C. and A. Kamranzadeh (2018). “The Optimal Peace Proposal When Peace Cannot Be Guaranteed”. *mimeo*.